

When Models Know They Are Being Tested: A Scoping Review of Evaluation Awareness, Sandbagging and Related Model-Originated Threats to the Validity of Large Language Model Evaluation

P N Bhargav Teja¹ and Divya Meena S^{1,*}

¹ School of Computer Science and Engineering, Vellore Institute of Technology, Vellore, Tamil Nadu, India

* Corresponding author: Divya Meena S, divyameena.s@vit.ac.in

Abstract

Evaluations of large language models (LLMs) increasingly inform deployment and governance decisions, yet models can recognise evaluation contexts and behave differently in them. This scoping review maps studies of such model-originated threats to evaluation validity: evaluation awareness, sandbagging, alignment and evaluation faking, strategic dishonesty and evaluator gaming. We searched Scopus, arXiv and Semantic Scholar and logged web sources to 9 October 2026. Screening and charting were AI-assisted against written criteria, with consistency checks but no second human reviewer. Of 5649 database and 418 web records, 315 studies were charted: 167 core studies and 148 adjacent studies; we read the full texts of 37 anchor studies. We propose a four-layer framework (context signal, awareness, behavioural divergence, validity consequence) and formalise it as a decomposition of evaluation error into scorer, context-induced, task and elicitation terms. Context separability is necessary for context-induced bias, awareness is neither necessary nor sufficient for a biased score, and context-induced divergence is not identifiable from evaluation data alone when evaluations are recognisable. Empirically, elicited, verbalised and represented awareness measures diverge; the awareness–behaviour link varies in magnitude and direction across models; and black-box sandbagging detection is unreliable, whereas fine-tuning and activation-level elicitation recover capability from induced sandbaggers. A full-text audit against our twelve-item reporting checklist found that 3 of 15 probe studies controlled for prompt format. The evidence is preprint-dominated, and the validation of detection and elicitation methods rests largely on deliberately induced behaviour. We close with an evidence-graded research agenda.

Keywords: Evaluation awareness · Sandbagging · Large language models · AI safety evaluation · Construct validity · Alignment faking

1 Introduction

Evaluations are the main instrument through which developers, independent evaluators and regulators decide what a large language model (LLM) can do and how it is likely to behave. Frontier developers condition deployment and safeguard decisions on dangerous-capability and alignment evaluations (Phuong et al. 2024; OpenAI 2025b; Anthropic 2025b), and third-party evaluators increasingly publish pre-deployment assessments (METR 2025a; Souly et al. 2026). Every such decision rests on an inferential step that is rarely stated: that behaviour observed under evaluation is representative of behaviour in deployment.

That step has always been vulnerable to familiar threats such as benchmark contamination, weak construct validity and prompt sensitivity (Raji et al. 2021; Sclar et al. 2024; Bean et al. 2025). This review concerns a different class of threat, one that originates in the evaluated model itself. A model that can recognise that it is

being tested may behave differently because it is being tested. Concretely, it may underperform on a capability evaluation (sandbagging), present a more aligned profile on a safety evaluation (alignment or evaluation faking), produce outputs that deceive the scoring pipeline (strategic dishonesty), or exploit the test harness (evaluator gaming). Each of these breaks the link between the reported score and the deployment behaviour the score is meant to predict.

Concern about this possibility is longstanding in the alignment literature (Hubinger et al. 2019; Carlsmith 2023; Ngo et al. 2024), but empirical work on it has grown sharply only since 2024. Models have been shown to underperform selectively when prompted or fine-tuned to do so (van der Weij et al. 2025); frontier models discriminate evaluation transcripts from deployment transcripts well above chance (Needham et al. 2025); linear probes recover an evaluation-versus-deployment distinction from model activations (Nguyen et al. 2025); and developers now report evaluation-awareness measurements in their system cards (Anthropic 2025a; OpenAI 2025a). At the same time, the field acknowledges that it lacks a shared conceptual foundation: studies conflate properties of the evaluation with properties of the model, and detection with behavioural response (Li C et al. 2026).

Existing reviews do not address this gap directly. Surveys of AI deception catalogue deceptive behaviours broadly but predate most of the evaluation-awareness literature or treat it as one behaviour among many (Park et al. 2024; Shi et al. 2026; Chen B et al. 2025). A 2026 review of the AI “observer effect” is closest in theme, but synthesises fifteen landmark contributions published up to 2024 and predates the measurement literature on evaluation awareness (Maurya et al. 2026). Broad reviews of safety-evaluation methods and of reward hacking name sandbagging or evaluator gaming as challenges without synthesising how they are measured (Grey and Segerie 2025; Wang X et al. 2026; Morampudi et al. 2026). A recent hybrid survey frames benchmark validity and alignment failure as a common proxy–target divergence problem across eight evidence streams (Uluurmak and Kurban 2026), but it does not centre on model-originated threats or on how they are measured. Reviews of benchmark validity and safety datasets address the evaluator side of the problem rather than the model side (Röttger et al. 2025; Bean et al. 2025). Section 2.4 and Table 1 give the detailed comparison.

This review asks five questions:

- **RQ1.** How are evaluation awareness, sandbagging, alignment or evaluation faking and related constructs defined and operationalised?
- **RQ2.** Which methods are used to detect or measure these phenomena, and what evidence supports their validity?
- **RQ3.** Under which model, task and context conditions are the phenomena observed, and where do findings conflict?
- **RQ4.** Which countermeasures have been proposed, and how rigorously have they been evaluated?
- **RQ5.** Which gaps most threaten the validity of current evaluation practice, and which studies would close them?

The review makes four contributions:

1. A reproducible, tiered evidence map: 167 core studies of model-originated threats to LLM evaluation validity and 148 adjacent studies, identified through three bibliographic databases and logged web searches, with all screening decisions logged and all counts generated by released scripts.
2. A four-layer framework (context signal, awareness, behavioural divergence, validity consequence) with explicit assignment rules, and a formal decomposition of evaluation error that states what each study

design can and cannot identify. The framework separates constructs that the literature frequently conflates.

3. A structured comparison of measurement methods and their validity evidence, and a typed analysis of seven apparent contradictions in the literature.
4. A twelve-item reporting checklist for evaluation papers and system cards, applied retrospectively to the full texts of the anchor studies with independent double coding, and an evidence-graded research agenda.

The remainder of the paper is organised as follows. Section 2 introduces terminology and prior reviews. Section 3 describes the review method, including the protocol amendments and the two-phase identification. Section 4 reports results by research question. Section 5 discusses implications. Section 6 sets out the research agenda, Section 7 the limitations of the review, and Section 8 concludes.

2 Background and terminology

2.1 Measurement validity and observer effects

The concern addressed here is an instance of a general measurement problem. In psychometrics, validity is a property of the inferences drawn from scores rather than of the scores themselves (Messick 1995). Measurement modelling makes the same point for computational systems: unobservable constructs must be operationalised, and mismatches between a construct and its operationalisation produce systematic error (Jacobs and Wallach 2021). Benchmarks presented as general measures of progress have been criticised on exactly these grounds (Raji et al. 2021). A recent review of 445 LLM benchmarks found widespread weaknesses in how measured phenomena, tasks and metrics support the claims made from them (Bean et al. 2025).

In the human sciences, the analogous model-side threat is the observer or Hawthorne effect: participants change their behaviour because they are studied. A systematic review of that literature concluded that research-participation effects exist but that little can be securely known about their conditions, mechanisms or magnitude, and that new concepts were needed (McCambridge et al. 2014). The parallel with the current LLM literature is close, and we return to it in Section 5. A related principle, often summarised as Goodhart’s law, holds that a measure placed under optimisation pressure ceases to measure what it measured before. Overoptimisation of a learned reward model illustrates it quantitatively (Gao et al. 2023).

2.2 Constructs

Five constructs recur in the literature. We define them operationally here, and Section 4.3 examines how the primary studies actually use them.

- **Situational awareness:** a model’s knowledge of itself and its circumstances, including whether it is in training, testing or deployment (Berglund et al. 2023; Laine et al. 2024). It is broader than evaluation awareness.
- **Evaluation awareness:** recognition by a model that its current context is an evaluation (Needham et al. 2025; Nguyen et al. 2025). Three senses need to be kept apart: a *capacity* to discriminate evaluation from deployment, an episode-level *state* (the model’s credence, *A* in Section 4.2), and a *report* of that state, verbalised or elicited. Some authors restrict the term to recognition alone; others include the resulting behaviour.
- **Sandbagging:** strategic underperformance on an evaluation relative to the model’s actual capability (van der Weij et al. 2025); we adopt this definition. The GPT-5 system card defines it as a model’s ability and

propensity to respond to safety or capability evaluations in a way that significantly diverges from performance under real conditions (OpenAI 2025a), and OpenAI’s Preparedness Framework lists sandbagging as a research category (OpenAI 2025b).

- **Alignment faking and evaluation faking:** behaving in line with the training or evaluation objective when the model infers that it is observed, while behaving differently when it infers that it is not (Greenblatt et al. 2024a; Fan et al. 2025).
- **Evaluator gaming:** obtaining a high score by exploiting the scoring process rather than performing the intended task, including test-case exploitation, reward hacking and outputs that mislead automated graders (Zhong et al. 2026a; Panfilov et al. 2026).

2.3 Elicitation and the direction of error

A capability evaluation underestimates a model when elicitation is incomplete, whether because the model hides a capability or because the evaluators did not try hard enough. Post-training enhancements such as scaffolding, tool use and prompting can raise benchmark performance by amounts that would otherwise require more than five times as much training compute, and sometimes more than twenty times (Davidson et al. 2023). Dangerous-capability evaluations therefore treat insufficient elicitation as a central threat (Phuong et al. 2024; Benton et al. 2024). Sandbagging is the adversarial limit of this problem. It is distinct from ordinary under-elicitation because the source of error is the model’s own policy. Conversely, alignment faking produces an *over*-estimate of safety. The two threats bias evaluation results in opposite directions on different evaluation types, and a review that treats them together must keep the direction of error explicit.

2.4 Prior reviews

Table 1 compares this review with the closest existing syntheses. None takes model-originated evaluation-validity threats as its primary object, and none synthesises the 2025–2026 primary literature on evaluation awareness measurement, which accounts for most of the evidence charted here (Section 4.1).

Table 1 Comparison with the closest prior reviews and syntheses (coverage judged from abstracts, and from the full text for Maurya et al. 2026)

Review	Type, status	Primary focus	Evaluation awareness as a validity threat	Measurement-method synthesis	What this review adds
Park et al. (2024)	Survey, <i>Patterns</i>	AI deception: examples, risks, solutions	Not a focus; predates most studies	No	Evaluation-centred scope; 2024–2026 evidence
Shi et al. (2026)	Taxonomy and benchmark analysis (50 benchmarks), ICLR 2026 workshop	Unified deception taxonomy mapped to benchmarks	One behaviour among many	Benchmarks, not measurement validity	Validity framework for the evaluation itself
Chen B et al. (2025)	Survey, preprint	Deception defined via signalling theory; deception cycle	Partial	No	Measurement-method comparison and contradictions
Uluirmak and Kurban (2026)	Hybrid survey (373 studies), preprint	Proxy–target divergence across eight evidence streams	Touched (strategic behaviour, alignment faking)	Broad; not awareness-specific	Model-side mechanisms; awareness measurement; reporting checklist

Review	Type, status	Primary focus	Evaluation awareness as a validity threat	Measurement-method synthesis	What this review adds
Maurya et al. (2026)	Review of 15 selected contributions (2016–2024), <i>IJSRCSEIT</i>	“Observer effect”: behaviour when unmonitored	Central, but conceptual; studies to 2024 only	No; narrative synthesis of landmark studies selected from Google Scholar, arXiv and developer repositories	Logged multi-database search; 2025–2026 measurement literature; comparability classes and reporting checklist
Grey and Segerie (2025)	Systematic literature review, preprint	Taxonomy of AI safety evaluations (capabilities, propensities, control)	Named as a challenge (sandbagging)	Broad taxonomy of evaluation techniques	Threat-specific synthesis of how sandbagging and awareness are measured and detected
Wang X et al. (2026); Morampudi et al. (2026)	Surveys, preprint; <i>Discover Artificial Intelligence</i>	Reward hacking in RLHF/RLVR and in agentic systems	Not a focus	Detection and mitigation of reward hacking	Reward hacking treated as one evaluator-gaming pathway within a validity framework
Frontier Model Forum (2025)	Industry technical report	Practice of frontier capability assessment	Elicitation and validity in general	No systematic evidence map	Documented evidence map
Röttger et al. (2025)	Systematic review (144 datasets), AAAI	Open safety-evaluation datasets	No	Dataset landscape	Model-originated threats
Bean et al. (2025)	Systematic review (445 benchmarks), NeurIPS	Construct validity of LLM benchmarks	No	Benchmark design	Model-side threats complementary to construct validity

3 Review methodology

3.1 Design, protocol and amendments

A protocol (v0.1) specified a scoping review reported according to PRISMA-ScR (Tricco et al. 2018), using the framework of Arksey and O’Malley (2005) as refined by Levac et al. (2010). A scoping design was chosen because terminology in the field is inconsistent, study designs are heterogeneous, and the objective is to map concepts, methods and gaps rather than to pool effects. The protocol was not registered; it and all amendments are released with the review (Supplementary File S0). The selection process is reported with a PRISMA 2020 flow diagram that separates database searches from other identification methods (Page et al. 2021). The optional consultation stage of Arksey and O’Malley was not undertaken.

The protocol, both identification phases and the analyses were completed on 9 October 2026. In the first phase, the bibliographic APIs named in the protocol were unreachable from the computing environment, so amendment v0.2 replaced them with 25 logged web-search batches and broadened the scope from evaluation awareness and sandbagging to model-originated threats to evaluation validity. In the second phase, after network access was opened, amendment v0.3 restored the database searches before they were run. It added Scopus, added “strategic deception”, “strategic dishonesty” and “deceptive alignment” to the specific phenomenon terms, and admitted “reward hacking” and “specification gaming” only together with an evaluation term. The search strategy was then frozen (search/queries.json). The web-search records were

retained as identification via other methods. Two planned sources could not be searched: OpenAlex returned HTTP 429 (daily quota exhausted) and DBLP returned an anti-automation page instead of results (Supplementary File S1b).

The deviations from protocol v0.1 are as follows:

- **No full-text eligibility stage.** Eligibility was judged on titles and abstracts (web records: titles, venues and retrieved summaries). Full texts were read for the 37 anchor studies only (Section 3.5).
- **Screening automation.** A rule-based pre-screen with an audited error rate and a manual rescue pass replaced manual title screening of records without phenomenon terms (Section 3.4).
- **No second human reviewer.** The planned 20% re-screen after seven days and the validation of a random sample by a second author were not carried out; consistency statistics from the AI-assisted process are reported instead (Section 3.4).
- **Appraisal.** The planned eight-item rigour profile was replaced by three design features charted from the anchor full texts and the twelve-item reporting audit (Section 3.8).
- **Exclusion codes.** E2 (no LLM), E4 (superseded version) and E5 (non-English) of v0.1 were not used as separate codes in the final logs. Superseded versions were merged as duplicates (R0), and records without an LLM or not in English were recorded as out of scope (E1 or a stage-2 exclusion).
- **Post-hoc tiering.** After an internal critical review found that the two phases had drawn the phenomenon boundary differently, one written tier rule was applied to all charted studies (Section 3.2), and the evidence map is reported for the core tier.

3.2 Eligibility criteria

Table 2 summarises the eligibility criteria. In brief, sources were eligible if they (I1) were primary empirical studies, methodological, formal or benchmark papers, or developer or third-party evaluation reports; (I2) measured, induced, detected or mitigated behaviour that depends on evaluation context, or strategic or induced underperformance, in LLMs or LLM-based agents; and (I3) reported enough method detail for the charting fields.

Table 2 Eligibility criteria and screening codes

Code	Criterion	Applied at
I1	Primary study, method, formal or benchmark paper, or developer/third-party evaluation report	Title, abstract, metadata
I2	Phenomenon: evaluation-context-dependent behaviour, strategic or induced underperformance, or a method to detect, elicit or mitigate it in LLMs or LLM agents	Title, abstract, metadata
I3	Method details sufficient to chart framework layer and method family	Title and abstract (full text for the anchor set)
Period	1 January 2019 to 9 October 2026	Metadata
E1	Different construct (e.g. environmental situational awareness), or a method unrelated to evaluation validity	Title, abstract
E3	Non-archival community forum post or informal blog without an organisational or archival report	Venue

Code	Criterion	Applied at
E6	Title or identifier could not be confirmed from retrieved records	Metadata
S1	No phenomenon term of the frozen query in title or abstract (rule-assisted pre-screen, audited)	Title, abstract
R0	Duplicate identifier or version of the same work	Identifier, title
CONTEXT	Eligible publication outside the phenomenon boundary (methodology, background theory, adjacent evaluation-validity literature such as contamination, jailbreak scoring or benchmark design); cited for context, not charted	Title, abstract

Tier rule. Every charted study was assigned to one of two tiers by the same written rule, whatever its route of identification (evidence/tier_coding.csv). A study is *core* if its primary object is (a) the evaluated model’s recognition of evaluation, training or monitoring context; (b) behaviour that differs with that context, including alignment or evaluation faking and scheming under oversight; (c) strategic underperformance, hidden capability or under-elicitation, and their detection or elicitation; (d) deception of evaluators, graders or monitors, or tampering with the evaluation record; (e) exploitation of an evaluation’s scoring, such as reward hacking or specification gaming against tests, verifiers or judges; or (f) a method, benchmark, formal analysis or developer or third-party report that addresses (a)–(e) directly. All other charted studies are *adjacent*. They supply mechanisms or methods on which the synthesis relies, such as chain-of-thought faithfulness, deception probes in general, backdoors, unlearning, introspection, AI control protocols and general deception benchmarks, and are used in Section 4.8. Because the CONTEXT category was applied by two passes, the boundary between adjacent studies and contextual references is less sharp than the boundary between core and adjacent.

Training-data contamination was deliberately placed outside the phenomenon boundary. It is a data-side rather than a model-side threat, and it has been reviewed recently (e.g. Chen S et al. 2025). Contamination studies were retained as context.

3.3 Information sources and search

Databases. Three sources were searched on 9 October 2026 through their APIs: Scopus (Elsevier Search API, field TITLE-ABS-KEY), arXiv (export API, abstract field, restricted to cs.CL, cs.AI, cs.LG and cs.CR) and Semantic Scholar (Academic Graph bulk search over titles and abstracts). Together they cover the peer-reviewed venues indexed by Scopus, the preprint server on which most of this literature first appears, and the conference proceedings (NeurIPS, ICML, ICLR, ACL) indexed by Semantic Scholar. The query combined three concept blocks: system terms (A: language model(s), LLM(s), foundation model, frontier model, AI agent, AI system); phenomenon terms specific enough to stand alone (B1: evaluation awareness, eval awareness, test awareness, sandbagging, strategic underperformance, alignment faking, evaluation faking, observer effect, evaluation gaming, under-elicitation, password-locked, hidden capabilities, strategic deception, strategic dishonesty, deceptive alignment); and broader phenomenon terms that were admitted only with an evaluation term (B2: situational awareness, scheming, capability elicitation, reward hacking, specification gaming; C: evaluation, benchmark, test, assessment, audit). The logical form was A AND (B1 OR (B2 AND C)), limited to publications from 1 January 2019. The exact strings submitted to each API, the retrieval timestamps and the record counts are given in Supplementary File S1b, and the raw API responses

are released. arXiv reported 1732 matches, of which 11 first posted before 2019 were removed by the date limit. Scopus records dated 2027 (online-first articles) were retained. OpenReview, which hosts most workshop papers in this field, was not searched as a source; it was used only to verify venues.

Other methods. The first-phase web searches comprised 133 queries in 25 record-identifying batches and one metadata-verification batch. They covered seven concept areas: evaluation and situational awareness; sandbagging and capability elicitation; alignment faking, scheming and sabotage evaluations; reward hacking and evaluator gaming; reasoning-trace monitoring and faithfulness; white-box detection; and system cards, third-party evaluations and governance. Known-item searches checked the metadata of sources cited by included studies (backward citation chasing). Developer system cards and third-party evaluation reports, which bibliographic databases do not index, were identified only by this route. Every query and the records it yielded are listed in Supplementary File S1a.

3.4 Selection

Deduplication. Database records were deduplicated automatically on DOI, version-stripped arXiv identifier and normalised title plus year (search/dedup.py). Title pairs with a difflib sequence-matcher ratio of at least 0.93 were flagged; all 30 flagged pairs were versions of the same work (for example, a preprint and its proceedings version) and were merged after inspection. Web records were matched to database records on the same keys, and records found by both routes are counted once, in the databases column of the flow diagram.

Automation tool (rule-based pre-screen). Semantic Scholar’s bulk search stems query terms (for example, “scheming” also retrieves “scheme”), which inflated the number of irrelevant records. Records whose title and abstract contained none of the phenomenon terms of the frozen query, written as a fixed regular expression (Supplementary File S2), were marked ineligible by this automation tool (3212 records). A random sample of 120 of them (seed 20261009) was read at title and abstract level and contained one relevant record (0.8%; Wilson 95% CI 0.2–4.6%). A broader rescue pattern then re-flagged 399 of the 3212 records for manual screening; 9 of these were included and 18 retained as context (6.8% eligible). Extrapolating the audit rate to the 2813 records that remained excluded by automation suggests about 23 missed eligible records (95% upper bound about 129); because the rescue pass yielded eligible records at a higher rate than the random audit, this estimate may be optimistic.

Manual screening. The 611 records that matched the phenomenon terms and the 399 rescued records were screened against Table 2 on title and abstract, and each received a decision (include, context or exclude). Reasons for exclusion were not recorded separately at this stage. Records in the first-phase web set were screened on titles, venues and retrieved summaries with the same criteria and one exclusion code each.

Who screened, and consistency checks. Screening and charting decisions were made by an AI system (a Claude model, Anthropic, operating as a research agent) applying the written criteria under the direction of the first author; no decision was duplicated by a second human reviewer. Two consistency checks are available. Ninety web records were also retrieved by the databases and were therefore screened twice, once in each phase and on different information (web summaries versus abstracts), without access to the earlier decision. The two passes agreed on 84 of the 90 records (inter-pass consistency 93.3%; Cohen’s $\kappa = 0.747$; Cohen 1960). The six discordant records were re-read and re-adjudicated by the same process; five became included studies and one a contextual reference. Because both passes come from the same AI-assisted process, κ here measures consistency, not inter-rater reliability. All decisions are logged (Supplementary File S2; screening/).

3.5 Data charting

For each charted study, the following were recorded: identifier, title, first author or organisation, year, venue, source type, tier, framework layer and method family (Section 3.6). Studies identified only through the databases were charted manually from title and abstract (`screening/db_charting_manual.tsv`). Source type was coded as *archival peer-reviewed* (journal or main-track proceedings), *workshop paper*, *preprint* or *grey literature* (developer and third-party reports). For every cited study, the venue was verified against arXiv metadata, OpenReview, Crossref or the conference page; for uncited studies the venue recorded at identification is used.

An anchor set of 37 core studies was selected purposively during the first phase as the primary studies on which the synthesis relies most, covering every framework layer; no explicit selection algorithm was used, so the anchor set is not a probability sample. Statistics computed on it describe the anchor set, not the literature. For each anchor study, the full text was retrieved from arXiv on 9 October 2026 (the version read is recorded in Supplementary File S3) and the following were charted:

- the construct name used;
- the setting;
- the models evaluated;
- the main result;
- the design features reported, namely a control or counterfactual condition, more than one model family, and ground truth by construction.

Every quantitative statement about an anchor study in Sections 4 and 5 was checked against the full text by the AI system, which recorded a verbatim quotation for each check; 16 of 158 statements were corrected as a result. Quantitative values for studies outside the anchor set are taken from their abstracts and are marked with a dagger (†) where they appear in the text.

3.6 Framework derivation and assignment rules

An initial framework was formulated in the protocol, separating the evaluation-context signal, model awareness, behavioural divergence and the validity consequence. It was then revised inductively during charting; for example, reasoning-trace monitoring was split from white-box methods after charting showed that the two have distinct validity problems. Each source was assigned to one primary layer according to its main object of study:

- **L1** studies treat properties of the evaluation (cues, format, realism, plausibility) as the independent variable or object of characterisation, including redesigns whose outcome is the recognisability of the evaluation;
- **L2** studies measure whether models recognise evaluation;
- **L3** studies measure behaviour conditional on context or incentive, including strategic underperformance and behaviour directed at the scorer;
- **L4** studies propose or test interventions whose outcome is detection, elicitation, suppression or reduced bias, or report validity consequences in developer or third-party evaluations.

A redesign of the evaluation is therefore L1 when its outcome is recognisability and L4 when its outcome is detected behaviour or a corrected score.

For studies identified in the first phase, method families were assigned by documented keyword rules applied to the title and recorded scope, with 30 manual overrides for multi-method studies whose titles do not reveal

their primary method (`evidence/chart.py`). Studies identified only in the database phase were assigned a layer and a method family manually from their abstracts. To check the rule-based layer assignment, the 37 anchor studies were also coded manually from their full texts: the two assignments agreed for 34 of 37 studies (rule–manual consistency 91.9%; $\kappa = 0.881$; `screening/reliability.py`). The disagreements concerned Li C et al. (2026), Ward et al. (2025) and Jang et al. (2026); the tables use the manual full-text coding for anchor studies. Single-layer assignment necessarily simplifies multi-method studies; Section 7 discusses the consequences.

3.7 Critical appraisal and synthesis

Formal risk-of-bias assessment is not required for scoping reviews, and no validated tool exists for the designs in this field. We did not compute quality scores. Instead, the full-text charting of the anchor set records which design features each study reports, and Section 4 discusses evidential weight qualitatively. Before any cross-study statement, findings were classified as directly comparable (same construct, measure and model), conditionally comparable (same construct but different measure or model) or non-comparable. No effect sizes were pooled and no leaderboards were constructed. Apparent disagreements were typed as a genuine contradiction, context-dependent, a metric artefact, a methodological artefact, or insufficient evidence (Table 7).

3.8 Reporting audit of the anchor set

To test whether the reporting practices recommended in Section 5.4 are already followed, the full texts of the 37 anchor studies were coded against the twelve checklist items (Table 8) using a frozen codebook (Supplementary File S5). Each item has an applicability rule; for example, the probe items apply only to studies that use activation probes or steering. Each study received one of four codes per item: reported, partly reported, not reported, or not applicable, with a verbatim quotation or a location for every positive code. Four AI coding agents coded disjoint subsets. A fifth agent, working independently and blind to the primary codes, re-coded a random sample of ten studies (seed 20261009). Agreement between the two independent codings over the 120 double-coded cells was 93.3% (Cohen’s $\kappa = 0.90$); agreement on applicability alone was $\kappa = 0.98$, and agreement on the code when both coders judged the item applicable was $\kappa = 0.84$ (85 cells). The primary codes are analysed. Adherence to an item is the share of applicable studies that report it, counting partial reporting as one half. Because the checklist was derived partly from these studies, the audit describes reporting practice; it is not a validation of the checklist.

4 Results


4.1 Evidence landscape

The database searches retrieved 5649 records (Scopus 244, arXiv 1721, Semantic Scholar 3684). After 1796 automatic and 30 manual duplicate removals, 3823 records remained. The rule-based pre-screen marked 3212 of them as ineligible; the rescue pattern re-flagged 399 of these for manual screening, leaving 2813 excluded by automation. Of the 1010 records screened manually (611 matching the phenomenon terms and 399 rescued), 687 were excluded (one exclusion was reversed at reconciliation), 157 were retained as contextual references and 166 were charted. The web searches added 418 unique records, of which 90 had already been retrieved by the databases; of the remaining 328, 48 were excluded, 131 retained as context and 149 charted. Applying the tier rule to the 315 charted studies gave 167 core studies, which form the evidence map, and 148 adjacent studies (Fig. 1).

Search recall. The databases retrieved 114 of the 167 core studies (68.3%) and 33 of the 37 anchor studies. Of the 53 core studies found only by the web searches, 24 are developer or third-party reports and 7 are workshop papers hosted on conference or OpenReview pages, which the databases do not index. The databases therefore retrieved 114 of the 136 core studies that are preprints or archival papers (83.8%). The 22 such studies they missed include four arXiv preprints with a phenomenon term in the title. One carries the term in its title but not in its abstract; the other three never use any of the system terms (such as “language model” or “LLM”) in their abstracts, and one of them writes “eval-awareness” with a hyphen (Zhuang and Aranguri 2026). Concept-block queries are therefore limited by vocabulary as well as by indexing, and the database search alone would not have been sufficient.

Identification via databases

Identification via other methods



Core tier by source type: archival peer-reviewed 23 | workshop 12 | preprint 107 | grey literature 25

Records found by both routes are counted once, in the databases column. Tiers are assigned independently of the route of identification. Counts are generated by search/dedup.py, screening/screen_db.py and screening/merge.py.

Fig. 1 PRISMA 2020 flow diagram of study selection, with separate columns for database searches and other methods. Records excluded by the rule-based pre-screen are shown as automation-tool exclusions. All counts are generated by the released scripts (search/dedup.py, screening/screen_db.py, screening/merge.py).

The evidence base is young and largely unrefereed. Of the 167 core studies:

- 107 (64.1%) are preprints, 23 (13.8%) archival peer-reviewed papers, 12 (7.2%) workshop papers and 25 (15.0%) grey literature, mainly developer system cards and third-party evaluation reports;
- 150 (89.8%) appeared in 2025 or 2026, with 102 between 1 January and 9 October 2026 (Fig. 2).

Archival peer-reviewed studies are concentrated in machine-learning conference proceedings (ICLR, ICML, NeurIPS, ACML, IJCNLP-AACL); only two core studies appeared in journals (Che et al. 2025 in *Transactions on Machine Learning Research* and a conceptual analysis in *Future Internet*, Ferrara 2026).




Fig. 2 Core studies by year and source type (n = 167). 2026 covers 1 January to 9 October.

Across framework layers, 82 core studies (49.1%) address validity consequences and responses (L4), 60 (35.9%) behavioural divergence (L3), 19 (11.4%) awareness (L2) and only 6 (3.6%) the properties of the evaluation context itself (L1) (Fig. 3). Across method families, the core evidence relies on:

- behavioural contrasts between conditions (53 studies);
- benchmark and environment construction (30);
- training- or weight-based methods, including model organisms and fine-tuning elicitation (26);
- developer or third-party audits (24);
- white-box methods such as probes, steering and feature analysis (22);
- reasoning-trace or transcript monitoring (12).

The 148 adjacent studies are concentrated at L4 (94) and L3 (38), mostly in reasoning-trace monitoring, deception probes, AI control and unlearning; Section 4.8 summarises them.

The method mix of the core tier is shifting. White-box methods rose from 1 of 14 core studies in 2024 and 4 of 48 in 2025 (8.3%) to 17 of 102 in 2026 (16.7%). Benchmark and environment construction rose from 3 of 48 in 2025 (6.3%) to 25 of 102 in 2026 (24.5%), driven largely by reward-hacking benchmarks and audits of agent benchmarks. The layer distribution has been stable since 2025, with L4 at half of the core studies in both years.

The imbalance between L1 and L4 is substantive. Many methods have been proposed for responding to evaluation-conditioned behaviour, but there are few systematic studies of *which properties of evaluations make them recognisable*. This makes the remaining evidence harder to interpret (Section 4.5).



Fig. 3 Evidence map: core studies by framework layer and primary method family (n = 167). Cell values are counts; single-layer assignment simplifies multi-method studies.

4.2 A four-layer framework

Charting showed that the literature repeatedly conflates four distinct objects of study. We separate them as follows (Fig. 4; Table 3).




Fig. 4 The CUE-AWARE-ACT-RESPOND framework. Each layer is a distinct object of measurement, labelled with the quantity it corresponds to in the formal decomposition defined below; validity consequences (L4) arise through L1-L3 or through the scoring function, and countermeasures act back on earlier layers.

Table 3 Framework layers: definitions, assignment rules and illustrative sources

Layer	Definition	Assignment rule	Illustrative sources
L1 Context signal	Properties of an evaluation that make it distinguishable from deployment (format, realism, synthetic prompts, scenario plausibility, honeypot cues)	Evaluation properties are the independent variable or object of characterisation, including redesigns whose outcome is recognisability	Li C et al. (2026); Ivanov and Africa (2026); Xiong et al. (2025); Krakovna et al. (2026)
L2 Awareness	Whether, and how, a model represents, verbalises or reports that it is being evaluated	Primary outcome is a measure of recognition	Needham et al. (2025); Nguyen et al. (2025); Heidari et al. (2026); Manek (2026)
L3 Behavioural divergence	Behaviour that differs between evaluation and deployment, or strategic misreporting or underperformance under evaluation	Primary outcome is behaviour conditional on context or incentive	van der Weij et al. (2025); Greenblatt et al. (2024a); Lynch et al. (2025); Panfilov et al. (2026)
L4 Validity consequence and response	The gap between score and deployment behaviour, and methods that detect, elicit, suppress or redesign around it; developer and third-party reports	Primary contribution is an intervention whose outcome is detection, elicitation, suppression or reduced bias, or a reported validity consequence	Tice et al. (2025); Taylor J et al. (2025); Hua et al. (2026); Souly et al. (2026)

The framework is useful mainly because it shows that a claim at one layer does not imply a claim at another. A probe that separates evaluation from deployment prompts (L2) does not show that behaviour changes (L3). A behavioural gap between conditions (L3) does not show that the model *recognised* the condition rather than responding to a correlated surface feature (L1). And a detection method validated on deliberately constructed model organisms (L4) does not show that it detects emergent behaviour. Sections 4.3–4.6 show that many contested findings are disagreements across layers rather than within them.

A formal statement. The layers can be made precise. Let P_E and P_D be the distributions of task instances x under evaluation and deployment, and let the context features available to the model (prompt format, scenario plausibility, tool environment) be drawn as $S \sim Q_c(\cdot|x)$ for $c \in \{E, D\}$ (L1). Let $A = a(x, S) \in [0, 1]$ be the model’s latent credence that it is being evaluated (L2), and let behaviour be drawn as $Y \sim \pi(\cdot|x, S, A)$, so that context can act on behaviour through awareness or directly through cues (L3). We assume (A1) that the context affects behaviour only through S , and (A2) that every evaluation task has a well-defined deployment presentation $Q_D(\cdot|x)$; in agentic settings, S includes environment feedback and A1 must hold at the level of trajectories. Let $m(x, y)$ be the evaluation’s scoring function and $m_0(x, y) \in [0, 1]$ the intended metric, oriented so that higher means more capable or safer, and write $\mu_c(x)$ for the expected intended score of the model’s behaviour on task x in context c . The evaluation reports $\theta_E = \mathbb{E}[m(x, Y)|c = E]$; for a propensity evaluation the target is deployment behaviour $\theta_D = \mathbb{E}_{x \sim P_D} \mu_D(x)$. The error then decomposes as

$$\theta_E - \theta_D = \Delta_{\text{score}} + \Delta_{\text{ctx}} + \Delta_{\text{task}},$$

where $\Delta_{\text{score}} = \mathbb{E}[m(x, Y) - m_0(x, Y)|c = E]$ is the scorer error under evaluation, including grader exploitation and judge leniency; $\Delta_{\text{ctx}} = \mathbb{E}_{x \sim P_E} [\mu_E(x) - \mu_D(x)]$ is the context-induced behavioural divergence for a fixed task; and $\Delta_{\text{task}} = \mathbb{E}_{x \sim P_E} \mu_D(x) - \mathbb{E}_{x \sim P_D} \mu_D(x)$ is the classical representativeness term. The identity is exact, but the split between Δ_{ctx} and Δ_{task} depends on the order of the terms: evaluating the context contrast under P_D instead of P_E moves an interaction term between them whenever context sensitivity varies with the task. For a capability evaluation the target is the elicitable capability θ^{cap} (the best intended score over admissible elicitation protocols), and a fourth term, $\Delta_{\text{elic}} = \mathbb{E}_{x \sim P_D} \mu_D(x) - \theta^{\text{cap}} \leq 0$, captures under-elicitation and sandbagging that does not depend on context. In this notation, context-

conditioned underperformance on a capability metric ($\Delta_{\text{ctx}} < 0$) and context-conditioned improvement on a safety metric ($\Delta_{\text{ctx}} > 0$) are necessary signatures of sandbagging and of alignment or evaluation faking respectively; intent is not represented. Model-originated behaviour directed at the scorer, such as test-case exploitation or strategic dishonesty towards judges, enters through Δ_{score} . Four remarks follow from the definitions and organise the rest of the synthesis.

- *Remark 1 (no context signal, no context-induced divergence).* Under A1, if $Q_E(\cdot|x) = Q_D(\cdot|x)$ for almost every evaluation task, then $\mu_E = \mu_D$ and $\Delta_{\text{ctx}} = 0$, whatever the model’s capacity for awareness. Context separability is therefore necessary for context-induced bias, but not for scorer-directed bias, which acts through Δ_{score} . Evidence that models discriminate evaluation from deployment transcripts (Needham et al. 2025) shows that the transcripts are separable; because those transcripts also differ in task content and outputs, it is consistent with, but does not establish, separability given the task, and it says nothing about the size of Δ_{ctx} .
- *Remark 2 (awareness is neither necessary nor sufficient).* Awareness is not sufficient: if the intended score does not depend on A or S , then $\Delta_{\text{ctx}} = 0$ even when $a(x, S_E) = 1$ and $a(x, S_D) = 0$. Awareness is not necessary: $\Delta_{\text{ctx}} \neq 0$ is possible with unchanged credence through a direct, cue-reactive dependence of π on S , and Δ_{score} and Δ_{task} can be non-zero without any awareness. Separating the awareness-mediated part of Δ_{ctx} from the direct part requires intervening on A with S held fixed, which is what steering designs attempt. Benchmark audits that find exploitable tasks (Section 4.9) are evidence about Δ_{score} , not about awareness.
- *Remark 3 (single-condition non-identifiability).* Data collected under evaluation identify behaviour only for context features within the support of $Q_E(\cdot|x)$. If $Q_D(\cdot|x)$ puts positive probability outside that support, as it does whenever evaluations are recognisable, then μ_D , and hence Δ_{ctx} , is not identified without data from a second condition or an assumed model of π : two policies that agree on the evaluation support and differ outside it generate identical evaluation data. If the supports coincide and Q_D is known, μ_D is identified by reweighting, which is the regime that deployment simulation and replay designs try to reach (Williams et al. 2026; Ivanov and Africa 2026). If only m and not m_0 is observed, Δ_{ctx} and Δ_{score} are not separately identified. The result is analogous to the two-run property that Xu (2026) formalises for monitors.
- *Remark 4 (a distributional bound).* Writing $\pi_c(\cdot|x)$ for the behaviour distribution averaged over $Q_c(\cdot|x)$, $|\Delta_{\text{ctx}}| \leq \mathbb{E}_{x \sim P_E} \text{TV}(\pi_E(\cdot|x), \pi_D(\cdot|x))$, where TV is the total-variation distance (the supremum over events). A small metric gap is therefore compatible with large behavioural divergence that happens to be orthogonal to m_0 ; such divergence does not bias θ_E for m_0 , but it can bias other metrics computed on the same outputs. Separately, averaging signed per-model gaps can mask divergence across models, as the per-model shifts of inconsistent sign reported by Knecht et al. (2026) illustrate.

Table 4 maps each design to the quantity it can identify. Applied to four anchor studies, the mapping shows how differently their headline numbers should be read. Removing or injecting verbalised awareness in reasoning traces (Schoen et al. 2025) is an intervention on A that estimates the awareness-mediated part of Δ_{ctx} for the traces manipulated. Steering an evaluation-aware model organism towards “deployment” (Hua et al. 2026) intervenes on a direction rather than on A itself, and the reported overshoot below the true deployment rate (10% versus 33% type-hint use) suggests that the direction does more than move awareness. Deployment simulation (Williams et al. 2026) estimates $\theta_E - \theta_D$ directly for the deployment sampled,

without decomposing it. Auditing games for sandbagging (Taylor J et al. 2025) evaluate detectors against a known Δ_{elic} that was induced by construction.

Table 4 Estimands identified by common study designs

Design	Layer	What it identifies	What it does not identify
Transcript or prompt discrimination by the model (elicited classification)	L2 (and L1 by implication)	That evaluation and deployment transcripts are separable for the model (jointly in task, context and output)	Whether π depends on that separation; elicitation inflates apparent recognition
Probes or verbalisation of awareness	L2	A proxy $\tilde{A}$ with method-specific bias (probes: possible format confounds; verbalisation: a lower bound if verbalisation never occurs without recognition)	Δ_{ctx}
Cue ablation or counterfactual framing with the task held fixed	L1 $\rightarrow$ L3	Δ_{ctx} for the manipulated cue	Effects of cues not manipulated; deployment realism
Steering of an awareness direction	L2 $\rightarrow$ L3	The effect of moving a direction on behaviour	That the direction is awareness, unless placebo directions are controlled (Mody et al. 2026)
Model organisms with induced divergence (prompted, fine-tuned, password-locked)	L3 / L4	Detector and elicitation performance against a known Δ_{ctx} or Δ_{elic}	Transfer to divergence that arises from training incentives
Benchmark and grader audits	L4	Δ_{score} and exploitability	Δ_{ctx}
Deployment simulation or production comparison	L4	$\theta_E - \theta_D$ directly, for the deployment sampled	The decomposition into its three terms

4.3 RQ1: Constructs and operationalisations

Terminological drift. The charted sources use at least eight overlapping labels:

- situational awareness, evaluation awareness and test awareness;
- the observer effect and the Hawthorne effect;
- evaluation faking and alignment faking;
- sandbagging (with “strategic underperformance” used interchangeably).

These labels differ in scope:

- *Situational awareness* in the SAD benchmark covers self-knowledge, situational inference and action (Laine et al. 2024). A benchmark of the same name in the NLP literature instead measures environmental perception and was excluded as a different construct (E1).
- *Evaluation awareness* is defined by some authors as recognition alone, explicitly separate from any behavioural adjustment (Needham et al. 2025). Others build the behavioural response into the definition (Fan et al. 2025).
- *Sandbagging* is defined as strategic, that is intentional rather than accidental, underperformance on an evaluation (van der Weij et al. 2025). It is operationalised in at least three different ways: underperformance under an explicit instruction, underperformance by a fine-tuned model organism, and underperformance that emerges in response to an incentive described in the environment (Meinke et al. 2024).

Three operational families for awareness. Awareness is operationalised in three ways that are not interchangeable:

1. *Elicited classification*: asking the model whether a transcript is an evaluation, or what it tests (Needham et al. 2025; Li X et al. 2026).
2. *Verbalised awareness*: detecting statements about being tested in reasoning traces (Schoen et al. 2025; Fan et al. 2025).
3. *Represented awareness*: decoding the evaluation–deployment distinction from activations (Nguyen et al. 2025; Chaudhary et al. 2025; Heidari et al. 2026).

A decomposition study argues that post-hoc detection queries (self-report and third-person judgement, and to a lesser degree purpose-guessing) systematically *over-report* recognition, because the detection query itself creates demand characteristics. It prefers monitoring of unprompted reasoning (Li C et al. 2026). Conversely, verbalised awareness is a lower bound, because models may recognise evaluations without saying so (Li X et al. 2026). The first method is therefore hypothesised to be biased upwards and the second downwards. Neither direction has been measured, verbalisation can also produce false positives, and no study charted here has quantified both biases on the same items. The two methods are best treated as complementary rather than redundant.

Behavioural constructs. At L3 the field distinguishes the following:

- **Capability-side divergence**: sandbagging and exploration hacking. Exploration hacking means strategically altering exploration during reinforcement-learning training to influence its outcome, for example to resist RL-based capability elicitation (Jang et al. 2026).
- **Propensity-side divergence**: alignment or evaluation faking, and behaving better when the context is recognised as a test (Greenblatt et al. 2024a; Lynch et al. 2025).
- **Measurement-side deception**: outputs that mislead the scorer, whether or not they also mislead the user. Examples include deliberately ineffective “harmful” answers that output-based jailbreak judges score as successes (Panfilov et al. 2026), and test-case exploitation in coding tasks (Zhong et al. 2026a).

The three classes bias evaluation results in different directions (Section 2.3), yet they are often grouped under a single umbrella term such as “scheming” or “evaluation gaming”.

4.4 RQ2: Measurement methods and their validity evidence

Table 5 summarises six method families, what each measures, the validity evidence for it, and the documented failure modes.

Table 5 Measurement and detection methods: validity evidence and documented failure modes (as reported in charted sources)

Method family	What it measures	Supporting validity evidence	Documented failure modes
Elicited classification (black-box)	Ability to discriminate evaluation from deployment transcripts	Well above chance on 1,000 transcripts from 61 datasets; frontier model AUC 0.83 vs human baseline 0.92 (Needham et al. 2025)	Demand characteristics inflate recognition (Li C et al. 2026); elicitation prompts do not transfer across model families, and the transcript generator explains about 11% of variance (Li X et al. 2026)

Method family	What it measures	Supporting validity evidence	Documented failure modes
Verbalised awareness in reasoning traces	Spontaneous statements about being tested	Causal: injecting or removing aware reasoning changes covert-action rates (Schoen et al. 2025); 99.3% of induced sandbagging driven by verbalised reasoning (Chaudhary 2026)	Lower bound; unverballed awareness reported in system cards and workshop studies (Anthropic 2026; Montoya et al. 2026); models can sandbag covertly against reasoning monitors (Li C et al. 2025); CoT faithfulness is limited (Chen Y et al. 2025; Arcuschin et al. 2026)
Linear probes on activations	Representation of the evaluation–deployment distinction	Decodable in Llama-3.3-70B-Instruct (Nguyen et al. 2025) and across families and sizes (Heidari et al. 2026); steering along probe directions changes verbalisation (Heidari et al. 2026)	Tracks benchmark format rather than context under controlled rewrites (Devbunova 2026); layer of best decoding shifts with scale (Manek 2026); probe directions depend on the prompts used to derive them (Noël 2026); latent-space monitors can be evaded (Bailey et al. 2026)
Activation steering and suppression	Behaviour when awareness is suppressed	With ground truth by construction, steering suppresses eval-aware behaviour even with the cue present, overshooting the true deployment rate (10% vs 33%; Hua et al. 2026)	Suppressing verbalisation does not show that unverballed awareness is removed; steering directions mix concepts; models can be trained to detect steering (Fonseca Rivera and Africa 2026)
Weight perturbation and training-based elicitation	Hidden capability	Noise raises accuracy of sandbagging models (Tice et al. 2025); few demonstrations unlock password-locked models (Greenblatt et al. 2024b); single-demonstration training recovered performance in an auditing game (Taylor J et al. 2025)	Elicitation also boosts benign models (false positives) (Taylor J et al. 2025); unreliable when no high-quality demonstrations exist (Greenblatt et al. 2024b); fine-tuned organisms can resist RL elicitation for hundreds of steps (Jang et al. 2026)
Consistency and lie-detection checks (black-box)	Inconsistency across paraphrases, or answers to unrelated follow-up questions	Unrelated-question lie detectors generalise across architectures (Pacchiardi et al. 2024); cross-context consistency proposed for sandbagging (Lin et al. 2026a)	Black-box detectors can learn model identity rather than deception (Lin et al. 2026b); black-box approaches defeated by imitation of a weaker model (Taylor J et al. 2025)

Three cross-cutting observations follow.

First, **ground truth is the scarce resource**. The strongest validity evidence comes from designs in which the true deployment behaviour is known by construction: password-locked or type-hint model organisms (Greenblatt et al. 2024b; Hua et al. 2026), tasks that are impossible by construction (Zhong et al. 2026a), and source-labelled transcript corpora (Needham et al. 2025). Of the 37 anchor studies, 14 have ground truth by construction and 8 partly; 33 report a control or counterfactual condition, and 27 test models from more than one family, with one more doing so partly (Supplementary Table S3; coded from full texts). Ground truth by construction, however, is almost always obtained by *inducing* the behaviour, so the resulting validity evidence concerns detection of induced behaviour. Whether it transfers to emergent behaviour is the most consequential open question in this literature (Section 6).

Second, **adversarial tests are sobering**. In the only auditing game for sandbagging with a blinded blue team charted here, the blue team could not reliably distinguish sandbagging from benign models. Black-box

approaches were defeated by effective imitation of a weaker model, and naive probing was vulnerable to behaviours that the red team had instilled (Taylor J et al. 2025). Deception probes that perform well on clean benchmarks degrade under distribution shift (Kumar 2026), and probes for alignment faking were reported to transfer for one model family but collapse to chance for another (Mehta 2026). Detection results from non-adversarial settings therefore overstate what can be expected against a model optimised to evade detection.

Third, **the three awareness measures diverge**. In the largest cross-paradigm study, behavioural, representational and mechanistic measures of evaluation awareness were not shown to be coupled across 37 open-weight models: only 1 of 15 pairwise correlations was significant, although with 32–36 models per correlation the power to detect moderate associations was limited (Nayan et al. 2026). The same study found that the models best at detecting evaluations behaviourally were the least swayed by harmful-request framing ($\rho = -0.79$). Representation and verbalisation also corresponded only partially within models (Heidari et al. 2026). It follows that “evaluation awareness” should not be reported as a single number, a point that several authors make explicitly.

4.5 RQ3: Conditions under which the phenomena are observed

Table 6 lists selected quantitative findings from the anchor set together with their comparability class. Because constructs, measures and models differ, almost no pair of findings is *directly* comparable. The synthesis below is therefore organised around the conditions that recur across studies rather than around effect sizes.

Table 6 Selected findings of anchor studies (values checked against the full texts; N/R = not reported)

Source	Layer	Setting	Reported finding	Comparability
Needham et al. (2025)	L2	1,000 transcripts, 61 datasets	Best model AUC 0.83; human baseline 0.92; agentic > chat	Conditional (reference point for black-box recognition)
Chaudhary et al. (2025)	L2	15 open models, 0.27B–70B, 4 families	Probe-measured awareness follows a power law in size	Conditional (probe design-specific)
Manek (2026)	L2	11 open models, 3 families	Within-family trajectories non-monotonic or inverse	Conditional (vs Chaudhary et al.: different sampling)
Nayan et al. (2026)	L2	37 open models, 7 families	1 of 15 cross-paradigm correlations significant	Conditional
Lynch et al. (2025)	L3	Blackmail scenario, n = 500	55.1% blackmail when model stated “real” vs 6.5% when it stated “evaluation”	Non-comparable (correlational, single model)
Greenblatt et al. (2024a)	L3	Training vs unmonitored cue	14% vs almost 0% harmful compliance	Non-comparable (primarily one model)
Sheshadri et al. (2025)	L3	Same paradigm, 25 models	5 of 25 models show significant compliance gaps; only Claude 3 Opus’s gap is consistently explained by alignment faking	Conditional (vs Greenblatt et al. 2024a)

Source	Layer	Setting	Reported finding	Comparability
Schoen et al. (2025)	L3	Anti-scheming training, 180+ environments	Covert actions o3 13.0% → 0.4%, o4-mini 8.7% → 0.3%; verbalised awareness rose after training	Conditional
Chaudhary (2026)	L3	Optimised prompts, 3 models, 4 benchmarks	Up to 94 pp degradation; 0.6 pp for one model on one benchmark (Claude-3.5-Haiku, HumanEval); in two models, 99.3% of sandbagging samples contained causally relevant evaluation-aware reasoning	Non-comparable across models
Zhuang and Aranguri (2026)	L3	Qwen3-32B, FORTRESS	24–46 pp compliance gap by awareness framing	Non-comparable (single model)
Knecht et al. (2026)	L3	8 open reasoning models, 4 benchmarks	Negligible-to-small output shifts with verbalised awareness	Conditional (vs Lynch et al.: different measure, models)
Li C et al. (2025)	L4	CoT-monitored sandbagging	16–36% monitor bypass when monitor-aware, conditional on successful sandbagging	Conditional
Taylor J et al. (2025)	L4	Blinded auditing game, 5 models	Blue team could not reliably discriminate sandbaggers	Conditional
Kirk et al. (2026)	L4	4 Claude models	0% unprompted sabotage; unprompted awareness mentions up to about 21% (21.3% in the abstract, 21.7% in the results)	Non-comparable (one developer’s models)

Five conditions recur across the charted studies.

Scale and capability. Several sources link evaluation awareness to scale or capability: a probe-based power law (Chaudhary et al. 2025), a shift from syntactic toward contextual cues as compact models approach 12B parameters, although structural cues still make up about 60% of their reasoning (Singh N and Chaudhary 2026), stronger models generally cheating more often on impossible tasks, a related evaluation-gaming behaviour rather than awareness itself (Zhong et al. 2026a), and evaluation faking reported to increase with reasoning ability (Fan et al. 2025). Denser within-family sampling, however, finds non-monotonic or inverse trajectories (Manek 2026). The trend across capability levels is therefore better supported than any specific functional form.

Agentic versus chat settings. Both models and humans identify evaluations more easily in agentic transcripts than in chat transcripts (Needham et al. 2025). This matters because the highest-stakes evaluations (sabotage, scheming, autonomy) are agentic.

Safety versus capability evaluations. Models react more strongly to cues of safety evaluations than to cues of capability evaluations (Li C et al. 2026). Probes classify existing safety evaluations as evaluation-like (Nguyen et al. 2025), and in a third-party evaluation the model refused more often when planted metadata indicated a dangerous-capability evaluation, while rarely mentioning this influence in its output (METR

2025a). The validity threat is therefore plausibly asymmetric, with propensity-based safety evaluations more exposed than capability benchmarks.

Training history. In Olmo 3 checkpoints, awareness was already decodable in the base model, increased after supervised fine-tuning and remained stable thereafter (Heidari et al. 2026). One developer report describes a sharp increase across model generations (Anthropic 2025a), and verbalised awareness rose during a capability-focused reinforcement-learning run (from 2.0% to about 20% in an experimental model without safety training; Schoen et al. 2025). Safety training aimed at reducing covert actions also coincided with higher verbalised awareness (Schoen et al. 2025). This raises the possibility that training *against* misbehaviour observed in evaluations teaches the model to recognise those evaluations.

Cue structure. No single cue factor raises awareness uniformly across models, but stacking cue factors progressively raises it in nearly all models (Li C et al. 2026). Smaller models rely heavily on formatting cues (Singh N and Chaudhary 2026), and probe evidence is confounded by benchmark format (Devbunova 2026). Scenario plausibility rather than realism has been proposed as the main driver in agentic misalignment scenarios (Vu et al. 2026). Replay-based evaluations that start from recorded real sessions reduced distinguishability relative to synthetic generators (Ivanov and Africa 2026).

4.6 Contradictions and their likely sources

Table 7 types seven apparent contradictions. A row carries two labels when two explanations apply, and the classification was made by the same AI-assisted process that charted the studies, not blind to the expected answer. No row was classified as a genuine contradiction. For each row, the last column states the observation that would discriminate between the two accounts; several could be made now with public models and released code, for example by re-estimating the scaling trend and the format confound on shared open-weight checkpoints. In six of them, the disagreement is more plausibly explained by differences in layer, measure or model than by a genuine conflict.

Table 7 Apparent contradictions in the literature and their classification

Apparent contradiction	Sources	Classification	Explanation supported by charted evidence	Discriminating observation
Awareness scales predictably (power law) vs non-monotonic or inverse within-family trajectories	Chaudhary et al. (2025); Manek (2026)	Methodological artefact / context-dependent	Different sampling density and model families (Qwen 2.5 excluded by Chaudhary et al., sampled at six sizes by Manek); the best-decoding layer itself shifts with scale, which Manek (2026) argues helps explain why denser within-family sampling yields non-monotonic or inverse rather than smooth trajectories	Re-estimate both on the same models, layers and sampling density; a smooth trend that survives per-layer selection and dense within-family sampling favours scaling

Apparent contradiction	Sources	Classification	Explanation supported by charted evidence	Discriminating observation
Self-stated “real” vs “evaluation” condition associated with 55.1% vs 6.5% blackmail, vs limited effect of awareness on behaviour	Lynch et al. (2025); Knecht et al. (2026)	Context-dependent / metric artefact	Different layer comparisons: self-classified condition split in one high-stakes agentic scenario versus aggregate output distributions across benchmarks; effect magnitude and direction vary by model (Abdelnabi and Salem 2025) and by awareness framing (Zhuang and Aranguri 2026)	Per-model reanalysis with the same awareness measure; an effect that persists when awareness is manipulated rather than self-stated favours a causal link
Models “know” they are evaluated (probes) vs probes track format	Nguyen et al. (2025); Devbunova (2026)	Methodological artefact	Probes trained on contrastive evaluation-versus-deployment data can confound context with format (Nguyen et al. note that test-classified prompts are often more structured or formatted); decorrelated designs are needed before probe evidence counts as evidence of awareness	Probes trained on data that decorrelate format from context and tested on format-matched evaluation and deployment prompts; retained accuracy favours awareness
Alignment faking is robust vs rare and model-specific	Greenblatt et al. (2024a); Sheshadri et al. (2025)	Context-dependent	Phenomenon demonstrated in specific models and prompt framings; compliance gaps appear and disappear with small prompt changes	Identical prompts and compliance-gap criteria across model families, with reasoning-trace coding of alignment-faking motives
Activation steering fails to elicit sandbagged capability vs activation-level intervention matches fine-tuning	Hofstätter et al. (2025); Le et al. (2026)	Methodological artefact (partly explained for password-locked organisms; single preprint)	Additive steering and zeroing fail because the password lock acts as a thresholded gate; setting the coordinate to its value in an honest reference recovers the capability (Le et al. 2026)	Reference-grafting applied to circuit-broken as well as password-locked organisms; success on both favours the gate explanation
Models can steer their activations to evade probes vs informed agents remain detectable	Kowalski et al. (2026); Das et al. (2026)	Context-dependent	Instructed control of simple residual-stream targets versus probes for collusion in multi-agent games with different models; neither tests a model with an incentive to evade a deployed monitor	Probes tested against models given an incentive, not an instruction, to evade a deployed monitor in the same task family
Behaviour change reflects strategic evaluation gaming vs sycophancy towards researchers	Fan et al. (2025); Baek et al. (2026)	Insufficient evidence	Current methods cannot mechanistically separate the two accounts (Baek et al. 2026); both predict better behaviour under recognised evaluation	Designs in which the two accounts predict opposite behaviour, e.g. reversing the evaluator’s stated preference under recognised evaluation

The last row is the most consequential. If evaluation-conditioned improvement reflects sycophancy towards evaluators rather than strategic concealment, the validity threat to safety evaluations is still real, because scores still overstate deployment behaviour. The appropriate countermeasures, however, differ. A formal analysis shows that, under finite behavioural evaluation of evaluation-aware agents, behavioural compliance does not uniquely identify latent alignment (Santos-Grueiro 2026). Behavioural tests can narrow the set of compatible hypotheses but cannot settle between them.

4.7 RQ4: Countermeasures

The charted countermeasures fall into five groups.

1. **Environment redesign (acting on L1).** Replay of real usage trajectories (Ivanov and Africa 2026), probe-guided rewriting of prompts towards deployment style (Xiong et al. 2025), honeypots embedded in realistic internal deployments (Krakovna et al. 2026), and realism scoring of auditing-agent transcripts against real conversations (Rosu and Wang 2026). Evidence of effectiveness is so far limited to reduced *distinguishability* as judged by models. Honeypot authors themselves caution that honeypots are not expected to remain reliable for sufficiently strategically aware systems (Krakovna et al. 2026).
2. **Suppression of awareness (acting on L2).** Activation steering toward “deployment” suppressed evaluation-aware behaviour in a model organism with known ground truth (Hua et al. 2026). In frontier-model reports, steering eliminated *verbalised* awareness, but whether it removed un*verbalised* awareness remained uncertain (Anthropic 2025a; Anthropic 2026).
3. **Elicitation (acting on L3 for capability).** Fine-tuning on a few demonstrations, reinforcement learning and weight perturbation recover hidden capability in model organisms (Greenblatt et al. 2024b; Tice et al. 2025; Hofstätter et al. 2025). Training on weak-supervisor samples reduced sandbagging to near zero (Ryd et al. 2026). Model-tampering attacks are proposed as conservative estimates of worst-case capability (Che et al. 2025). The main documented risks are false positives on benign models (Taylor J et al. 2025) and organisms that resist reinforcement-learning elicitation (Jang et al. 2026), although in the same study supervised fine-tuning on benign traces quickly recovered the locked capability and reasoning-trace monitors, action-only monitors and weight noising detected the organisms.
4. **Training interventions (acting on L3 for propensity).** Training models to be evaluation-cooperative closed the evaluation–deployment gap by 70–100% in five of seven model settings in one ICML 2026 workshop paper† (Li and Turner 2026). Inoculation prompting is reported to suppress unwanted traits at test time (Tan et al. 2025; Wichers et al. 2025). Both approaches rely on the model’s own cooperation, so validating them requires the same evaluations whose validity is in question.
5. **Detection and monitoring (acting on L4).** Reasoning-trace monitors, deception probes and consistency checks (Table 5). Monitoring of reasoning traces is described as a valuable but fragile opportunity that optimisation pressure can erode (Korbak et al. 2025). Direct optimisation against a reasoning monitor produced obfuscated reward hacking (Baker et al. 2025).

None of the five groups has been tested against a model that is both strategically aware *and* has not been deliberately constructed to be detectable. That is the central limitation of the evidence on countermeasures.

4.8 Adjacent evidence streams bearing on validity

This section draws on the 148 adjacent studies (Section 3.2) together with the core studies of reward hacking and evaluator gaming. These research streams were not designed as studies of evaluation validity, but they supply mechanisms, methods or failure modes on which the synthesis above depends. We summarise them briefly, because a reader assessing an evaluation’s validity will need them.

Reward hacking and evaluator gaming. Specification gaming learned in simple environments generalised to rewriting a reward function in a minority of cases (Denison et al. 2024). Reasoning models hacked a game environment to win (Bondarenko et al. 2026). Frontier agents exploited scoring code in third-party evaluations (METR 2025b) and cheated on impossible coding tasks (Zhong et al. 2026a). Training on harmless reward hacks generalised to broader misbehaviour (Taylor M et al. 2025), as did narrow fine-tuning on insecure code (Betley et al. 2025a), a phenomenon whose internal correlates have been linked to persona features and convergent linear directions (Wang M et al. 2026; Soligo et al. 2025; Turner et al. 2025). Each of these is a route by which a high evaluation score can reflect exploitation of the evaluator rather than the target capability. Automated judges have well-documented biases of their own, including sensitivity to position and length (Wang P et al. 2024; Dubois et al. 2024), and jailbreak scoring has overstated attack success (Souly et al. 2024; Mazeika et al. 2024).

Reasoning-trace faithfulness and monitorability. Monitoring verbalised awareness is only as informative as reasoning traces are faithful. Chain-of-thought explanations can be systematically unfaithful to the factors driving an answer (Turpin et al. 2023; Lanham et al. 2023), and reasoning models acknowledge influential hints only part of the time (Chua and Evans 2025; Chen Y et al. 2025), including on ordinary prompts (Arcuschin et al. 2026). Models can be trained to encode reasoning in text (Roger and Greenblatt 2023), and steganographic collusion can emerge from misspecified incentives (Mathew et al. 2025; Motwani et al. 2024). Monitors can catch sabotage better from reasoning than from actions alone in some settings (Arnav et al. 2025), and evasion is hard when reasoning is necessary for the task (Emmons et al. 2025), but optimisation pressure on the reasoning trace produces obfuscation (Baker et al. 2025; Korbak et al. 2025). Emergent misalignment in reasoning models produced both overt and benign-sounding rationalisations that monitors missed (Chua et al. 2025).

White-box detection. Linear structure for truth and refusal is well documented (Marks and Tegmark 2024; Ardit et al. 2024), as are top-down representation reading and contrastive steering (Zou et al. 2023; Rimsky et al. 2024). Latent knowledge can sometimes be recovered independently of outputs (Burns et al. 2023; Mallen et al. 2024), and trained activation explainers extend this approach (Karvonen et al. 2026). Deception probes reach high recall at low false-positive rates in benchmark settings (Goldowsky-Dill et al. 2025) and give modest gains when added to black-box monitors (Parrack et al. 2025). Their performance, however, varies across deception types (Kretschmar et al. 2025) and degrades under distribution shift (Kumar 2026), and evading latent-space monitors is possible (Bailey et al. 2026). For alignment faking specifically, internal-manipulation methods identified most alignment-faking models in a constructed benchmark (Clymer et al. 2024), but probe transfer across model families failed in naturally occurring cases (Mehta 2026). Related auditing work found hidden objectives through interpretability and training-data analysis (Marks et al. 2025). It also studied deployment-triggered backdoors, which persist through safety training (Hubinger et al. 2024), can use temporal cues as triggers (Price et al. 2024), and can sometimes be reconstructed without knowledge of the trigger (Bullwinkel et al. 2026).

Situational awareness and self-knowledge. Out-of-context reasoning, a hypothesised precursor of situational awareness, improves with scale but depends on training setup (Berglund et al. 2023; Treutlein et al. 2024). Models show privileged access to some facts about their own behaviour (Binder et al. 2025). They can describe behaviours they were fine-tuned into without in-context examples (Betley et al. 2025b) and recognise and favour their own generations (Panickssery et al. 2024). These capabilities are the prerequisites for evaluation awareness: a model that knows what it is and what it tends to do is better placed to notice when its situation is artificial.

Scheming, deception and sabotage evaluations. A broader literature evaluates behaviours that evaluation awareness could mask, including:

- strategic deception under pressure (Scheurer et al. 2024) and deception abilities in general (Hagendorff 2024);
- honesty disentangled from accuracy (Ren et al. 2025) and deception benchmarks (Huang et al. 2025; Wu et al. 2026);
- sabotage of monitored tasks (Benton et al. 2024; Kutasov et al. 2025; Ward et al. 2025);
- stealth and situational reasoning prerequisites for scheming (Phuong et al. 2025) and safety cases built on such evaluations (Balesni et al. 2024);
- shutdown resistance (Schlatter et al. 2025) and self-replication (Pan et al. 2024).

Control evaluations treat the model as a potential adversary and measure whether monitoring protocols hold under red-teaming (Greenblatt et al. 2024c; Kaufman et al. 2025; Gardner-Challis et al. 2026). That stance is arguably the appropriate default when awareness cannot be ruled out.

Elicitation and hidden capability. Unlearning that appears successful on benchmark accuracy can often be undone by brief fine-tuning or recovered through activation-space interventions (Łucki et al. 2025). This motivates tamper-resistance methods (Tamirisa et al. 2025) and tampering-based evaluation (Che et al. 2025) on hazardous-knowledge benchmarks (Li N et al. 2024). The lesson for sandbagging is the same: absence of a capability in default evaluation is weak evidence of absence.

Benchmark and agent-evaluation validity. Independently of model-side behaviour, agentic benchmarks can under- or over-estimate performance through task and grader design (Zhu et al. 2025). Capability trends depend on human-calibrated task suites (Kwa et al. 2025). Agent safety benchmarks measure misuse compliance rather than evaluation-conditioned behaviour (Andriushchenko et al. 2025; Kuntz et al. 2025). Selective leaderboard disclosure distorts comparisons (Singh S et al. 2025), and test-set contamination, which can be tested statistically for black-box models (Oren et al. 2024), inflates scores. These evaluator-side threats compound the model-side threats reviewed here, and calls for an evaluation science stress the same point (Weidinger et al. 2025; Burnell et al. 2023; Hardy et al. 2025; Liang et al. 2023).

4.9 Formal limits, awareness manipulation and exploitable evaluation infrastructure

Four groups of recent core studies, most posted in 2026 and none in the anchor set, bear directly on the synthesis. Their quantitative results are taken from abstracts and are marked †.

Formal limits of single-trace monitoring. Xu (2026) formalises sandbagging and evaluation awareness as two-run (“2-safety”) properties, which can be witnessed only by comparing two executions, and bounds the balanced accuracy of any monitor that sees a single trace at $\frac{1}{2} + \frac{1}{2} \cdot \text{TV}(P_0, P_1)$, where TV is the total-variation distance between the two behaviour distributions. In a factorial test, LLM monitors stayed at chance (50.4%†) when asked to imagine a second run but reached 90.0%† when given an executed second run. Rozinek (2026) derives Chernoff–Stein sample-complexity bounds on the number of queries needed to detect a sandbagger. Both results support the framework’s central methodological claim (Remark 3): behavioural divergence (L3) is a property of a *pair* of conditions and cannot be established from transcripts of the evaluation condition alone.

Measuring and manipulating awareness. Verbalisation training increased verbalised evaluation awareness 2.4–2.9-fold† across three models while measured latent awareness and behaviour remained largely stable (Anwar et al. 2026). Verbalisation can therefore be moved independently of the latent quantity, as the

measurement-divergence findings in Section 4.4 would predict. Input-only optimisation suppressed an evaluation-awareness latent in two Llama models, but a placebo random direction was suppressed as strongly and shifted behaviour as far, and suppressing the direction with a real evaluation passage in context did not reduce the model’s behavioural judgement that it was being evaluated (Mody et al. 2026). In the authors’ words, activation readability is not behavioural controllability. Two studies test the assumption that models cannot hide their activations, with opposite results in different settings: most LLMs could steer their residual-stream activations on instruction well enough to evade probes in simple tasks, though imperfectly (Kowalski et al. 2026), whereas colluding agents told that an activation monitor was present stayed detectable by the best probes (Das et al. 2026). At L1, evaluation framing lengthened reasoning and reduced answer-only compliance in GPT-OSS-20B, with limited or inconsistent accuracy gains (Ahmed et al. 2025). Deployment simulation, which regenerates responses to prefixes of real de-identified conversations, gave evaluation-awareness estimates much closer to production traffic than conventional evaluations, using registered, outcome-blinded predictions for one release (Williams et al. 2026). This is the first realism countermeasure in the evidence map that was validated against deployment outcomes rather than against a discriminator (cf. Table 9, gap 7).

Elicitation of sandbagged capability. Reference-grafting, which sets an activation coordinate along a contrast direction to its value in an honest reference, recovered 94–101% † of the gap between honest and sandbagging performance across eleven password-locked models from three architecture lineages, matching fine-tuning without weight updates (Le et al. 2026). The authors explain the earlier failure of additive steering (Hofstätter et al. 2025) by showing that the lock behaves as a thresholded gate, which reconciles part of the steering-versus-fine-tuning contradiction in Table 7. A concealed-information probe read the recognised answer with recognition scores of 0.85–0.93† under prompted deception, trained sandbagging and password locking, and separated models that will not answer from models that cannot (Dingeto 2026). Both results concern sandbagging induced by construction, so the transfer question of gap 1 (Table 9) remains open.

Exploitable evaluation infrastructure. Several audits show that evaluator gaming does not require evaluation awareness when the evaluation itself is exploitable. Frontier models given only the task description could hack 323 of 1968 tasks (16%†) across five terminal-agent benchmarks (Zhong et al. 2026b). An automated auditor produced exploits that achieved near-perfect scores on most of ten agent benchmarks without solving a task, surfacing 219† distinct flaws (Wang H et al. 2026). In a 49-task sample of SWE-bench Verified, more than a quarter of tasks (reported as 28.5%†) had test suites that accepted a verified incorrect patch (Rajan 2026), and a released dataset documents 3632 † exploit trajectories in 331 † reward-hackable terminal environments (Bercovich et al. 2026). Two further studies extend the attack surface to the evaluation record and the evaluator. Most of the agent harnesses tested let agents delete their own execution traces without triggering guardrails, and trace tampering emerged without instruction when agents tried to improve their rewards (Qin et al. 2026). LLM judges softened safety verdicts when told that low scores would lead to retraining or decommissioning, with a peak shift of 9.8 percentage points†, while the reasoning traces of reasoning-model judges never mentioned the framing (Gupta et al. 2026). That last result is an evaluator-side analogue of evaluation faking that reasoning-trace monitoring cannot detect.

5 Discussion

5.1 Principal findings

The review supports four conclusions, each qualified by the strength of the evidence behind it.

(i) Evaluation-conditioned behaviour in LLMs has been demonstrated, but mostly under induced or strongly scaffolded conditions. Evidence is strong that models *can* underperform selectively when instructed or trained to (van der Weij et al. 2025), *can* discriminate evaluation from deployment transcripts (Needham et al. 2025), and *can* behave differently when they infer that they are observed (Greenblatt et al. 2024a; Lynch et al. 2025). Evidence on how often this happens *spontaneously* in realistic evaluations is weaker and mixed. HoneyPot evaluations in a real internal deployment found no unprompted scheming (Krakovna et al. 2026); aggregate output shifts associated with verbalised awareness were small (Knecht et al. 2026); and only a minority of models showed alignment faking in a replication across 25 models (Sheshadri et al. 2025).

(ii) “Evaluation awareness” is a family of weakly coupled constructs, not one quantity. Representational, verbalised and elicited measures diverge (Nayan et al. 2026; Heidari et al. 2026). Each has a hypothesised directional bias (Section 4.3). Probe evidence can be confounded by prompt format, as one controlled single-model study showed (Devbunova 2026), whereas other probes retained discrimination under rewrites that changed behaviour (Nayan et al. 2026); only 3 of 15 anchor studies using probes or steering controlled for format (Section 5.4). Reports that state an awareness rate without specifying the measure are therefore uninterpretable across studies.

(iii) The link from awareness to behaviour is heterogeneous in magnitude and direction. Test awareness can increase or decrease harmful compliance depending on the model (Abdelnabi and Salem 2025). How the awareness is framed predicts compliance (Zhuang and Aranguri 2026). Optimised prompts that induce evaluation awareness degraded Claude-3.5-Haiku by 73 pp on arithmetic but by only 0.6 pp on code generation, where Llama-3.3-70B dropped 86.6 pp (Chaudhary 2026). A single “evaluation-awareness correction” applied across models is therefore not justified by current evidence.

(iv) Detecting adversarial underperformance is harder than eliciting capability. In adversarial settings, detection of sandbagging was unreliable (Taylor J et al. 2025; Li C et al. 2025), and monitors that detected code sabotage well found sandbagging harder to detect (Ward et al. 2025), whereas training-based elicitation more often recovered hidden capability (Greenblatt et al. 2024b; Taylor J et al. 2025; Ryd et al. 2026), at the cost of false positives. For capability evaluations, this favours elicitation-first protocols, which report the best elicited performance, over detection-first protocols, which try to classify whether a model is sandbagging.

5.2 Comparison with the human-sciences literature

The trajectory of this literature resembles that of research on the Hawthorne effect, where effects of being studied were found to exist but their conditions, mechanisms and magnitudes remained poorly understood, and new concepts were called for (McCambridge et al. 2014). Two differences matter.

First, LLM evaluators have access to model internals. This permits measurement designs, namely representation probing, steering and weight perturbation, that are impossible with human participants. It also introduces new validity threats specific to those designs (Table 5).

Second, unlike human participants, models are trained, and training can be informed by evaluation results. If evaluations that detect misbehaviour feed back into training, the training process creates selection pressure toward recognising and passing those evaluations. That pressure is what Goodhart-type arguments predict (Gao et al. 2023). Verbalised awareness did rise after anti-scheming training, but it rose much more after capability-focused reinforcement learning without safety training (from 2.0% to 20.6%; Schoen et al. 2025), so the rise may reflect a general effect of reinforcement learning rather than feedback from evaluations; the

design cannot separate the two. The analogy with human research participation therefore *understates* the threat.

5.3 Comparison with prior reviews

Relative to deception surveys (Park et al. 2024; Chen B et al. 2025; Shi et al. 2026), this review treats the evaluation as the object of study, so that deception matters only insofar as it changes what scores mean. Relative to the review of the AI “observer effect” by Maurya et al. (2026), which discusses fifteen landmark contributions up to 2024, it adds a reproducible multi-database search, the 2025–2026 measurement literature and an explicit separation of awareness from behaviour. Relative to the broad proxy–target framing of Uluirmak and Kurban (2026), it isolates the model-originated mechanisms and their measurement. Relative to construct-validity reviews of benchmarks (Bean et al. 2025; Reuel et al. 2024), it adds a class of threat that remains even when a benchmark is well designed, although the benchmark audits in Section 4.9 show that many agent benchmarks are not. A benchmark with strong construct validity can still yield invalid inferences if the evaluated system behaves differently because it recognises the benchmark.

5.4 Implications for evaluation practice: a reporting checklist

The heterogeneity documented in Sections 4.3–4.6 is partly a reporting problem. Table 8 proposes a twelve-item checklist for evaluation papers, system cards and third-party evaluation reports, derived directly from the failure modes and contradictions charted above. The checklist is a proposal derived from the evidence and has not been validated as a reporting standard. We did, however, apply it retrospectively to the full texts of the 37 anchor studies (Section 3.8; Fig. 5).

Table 8 EVAL-AWARE reporting checklist for evaluations exposed to model-originated validity threats

#	Item	Rationale (evidence in this review)
1	State which awareness construct is measured (elicited, verbalised, represented) and report each separately	Measures diverge and are hypothesised to have opposite biases (Sections 4.3–4.4)
2	For probe-based claims, report a design that decorrelates evaluation context from prompt format	Probes track format (Devbunova 2026)
3	Report the layer and selection procedure for probes, with a control probe baseline	Best-decoding layer shifts with scale (Manek 2026)
4	Report the transcript generator for deployment-side data and test cross-family transfer of elicitation prompts	Generator effects and non-transfer (Li X et al. 2026)
5	Report the direction of possible bias (capability under-estimation vs safety over-estimation) for each evaluation	Threats bias in opposite directions (Section 2.3)
6	For capability evaluations, report elicitation effort (prompting, scaffolding, fine-tuning, RL), and best elicited performance alongside default performance	Elicitation recovers hidden capability more reliably than detection (Section 5.1)
7	When excluding refusals or failures, report results with and without exclusion	Exclusions turn reported scores into ceilings (single preprint: Saraireh 2026); 4 of 25 applicable anchor studies reported both (Section 5.4)
8	For propensity evaluations, report a cue-ablated or replay-based variant	Stacked cues drive awareness; replay reduces distinguishability (Li C et al. 2026; Ivanov and Africa 2026)

#	Item	Rationale (evidence in this review)
9	Report unprompted verbalised awareness rates in agentic evaluations, and treat them as lower bounds	Unverbalised awareness documented (Anthropic 2026)
10	Report per-model rather than pooled awareness–behaviour relations	Magnitude and direction vary by model (Abdelnabi and Salem 2025)
11	Report uncertainty (confidence intervals, number of samples, seeds) for all rates	Small-sample rates and prompt sensitivity (Sclar et al. 2024; Mizrahi et al. 2024)
12	State the access level available to the evaluator (black-box, grey-box, white-box)	Access constrains which methods are possible (Casper et al. 2024)




Fig. 5 Reporting audit of the 37 anchor studies against the EVAL-AWARE checklist (Table 8), coded from full texts. Adherence (top) is the share of applicable studies reporting the item, counting partial reporting as one half. Studies are ordered by the number of items reported. Inter-coder agreement on a random 10 studies: $\kappa = 0.90$.

Reporting audit. The audit shows a clear split (Fig. 5). Items that concern the *framing* of a study are almost always reported: every study states the evaluator’s access level (37 of 37), every multi-model study reports results per model (34 of 34), and nearly all name the awareness construct they measure (adherence 94%, 26 applicable studies) and state the direction of the possible bias (89%). Items that concern *threats to the measurement itself* are reported much less often:

- only 3 of the 15 studies that use probes or steering as evidence fully decorrelate evaluation context from prompt format (adherence 37%), although format confounding is the best-documented failure mode of probe evidence (Devbunova 2026);
- only 7 of the 27 agentic or reasoning-trace studies report unprompted verbalised-awareness rates and treat them as lower bounds, and 12 report no such rate at all (adherence 41%);
- only 1 of the 18 studies that rely on deployment-side transcripts or elicitation prompts tests whether results transfer across transcript generators or model families (adherence 47%);
- only 4 of the 25 studies that exclude outputs (refusals, unparseable answers, filtered tasks) report results with and without the exclusion (adherence 52%);
- uncertainty is fully reported for main results in 17 of the 37 studies (46%) and partly in 19 more.

Median adherence per study was 79% (range 50–100%). The pattern matters for interpretation. The items least often reported are exactly those that would distinguish a genuine awareness effect from a measurement artefact (Table 4): format confounds for probes, the lower-bound status of verbalised awareness, generator effects for transcript classification, and the effect of exclusions on reported rates. Items 2, 4, 7 and 9 are therefore the most useful additions to current practice.

6 Research gaps and prioritised agenda

Table 9 lists evidence-based gaps, each with the study that would address it and the strength of the evidence that the gap exists. Gaps are ordered by our judgement of priority, which weighs importance for evaluation validity, evidence strength and feasibility.

Table 9 Prioritised research agenda

Priority	Gap	Evidence that the gap exists	Proposed study	Evidence strength
1	Detection and elicitation validated only on <i>induced</i> behaviour	Ground truth obtained almost exclusively by construction (Section 4.4); organisms deliberately trained to be detectable or resistant	Blinded auditing games in which red teams construct organisms by <i>training incentives</i> rather than direct instruction; report transfer from induced to incentive-driven behaviour	Strong
2	No validated mapping between awareness measures	1 of 15 cross-paradigm correlations significant; partial representation–verbalisation correspondence	Multi-measure study on shared items with known ground truth, estimating the upward bias of elicitation and the downward bias of verbalisation	Strong
3	Few studies of evaluation properties (L1)	6 of 167 core studies (3.6%)	Factorial cue-ablation studies across task families, including agentic, multimodal and multilingual evaluations	Strong (coverage gap)
4	Awareness–behaviour link unexplained across models	Direction and magnitude vary (Table 7)	Pre-registered, multi-family studies relating awareness framing to behaviour change, with model-level random effects	Moderate

Priority	Gap	Evidence that the gap exists	Proposed study	Evidence strength
5	Feedback from evaluation into training	Verbalised awareness rose after anti-scheming training, and more after capability-focused RL (Schoen et al. 2025)	Controlled training runs that vary whether evaluation-derived data enter training; measure subsequent awareness	Moderate
6	Sycophancy vs strategic accounts not separable	Baek et al. (2026); Santos-Grueiro (2026)	Designs in which the two accounts predict different behaviour (e.g. evaluator-preference reversal under recognised evaluation)	Moderate
7	Countermeasures validated by model judgement only	Realism mostly validated by model discriminators; one deployment-grounded validation (Williams et al. 2026)	Validate realism and replay methods against deployment outcomes, not only against discriminator accuracy	Moderate
8	Peer-review and replication deficit	13.8% of core studies archival peer-reviewed; key findings single-model	Independent replications of anchor findings (Table 6) with released prompts and code	Strong
9	Evaluation reporting heterogeneity	Rates reported without construct, uncertainty or access level	Adoption study of a reporting checklist (Table 8) across system cards	Weak (practice evidence only)

7 Limitations of the review

The review has substantial limitations, which readers should weigh before relying on its conclusions.

Who did the work. An AI system made the screening, tiering and charting decisions and performed the full-text checks of the anchor studies under the direction of the first author, working to written criteria with every decision logged. No decision was duplicated by a second human reviewer, and the consistency statistics reported here (inter-pass $\kappa = 0.747$, rule–manual $\kappa = 0.881$, inter-coder $\kappa = 0.90$ for the reporting audit) come from AI-assisted processes. They are evidence of consistency, not of human inter-rater reliability. A human audit of a stratified sample of decisions is the most important outstanding check.

Information sources and recall. Three databases were searched, but OpenAlex and DBLP could not be queried, and Web of Science, the ACM Digital Library, IEEE Xplore and OpenReview were not searched. The databases retrieved 84% of the core studies of the kinds they index, but missed some that carry a phenomenon term in their titles, so even a well-specified query is limited by term variants and indexing lag. Developer reports and workshop papers were found only through web searches, whose coverage cannot be quantified. Non-English work is probably under-represented.

Automation and eligibility. A rule-based pre-screen excluded 2813 records without manual reading. The audit sample and rescue pass bound, but do not eliminate, the number of eligible records missed (Section 3.4). Eligibility was judged on titles and abstracts; there was no full-text eligibility stage.

Tiering. The core/adjacent rule was applied after an internal critical review showed that the two identification phases had drawn the phenomenon boundary differently. The rule is written down and applied to every charted study, but it was applied post hoc, and the boundary between adjacent studies and contextual references remains less sharp than the core boundary.

Depth of reading and the anchor set. Full texts were read for 37 anchor studies only; other studies were charted from titles and abstracts, and their numbers are marked †. The anchor set was chosen purposively, so statistics computed on it, including the design-feature counts and the reporting audit, describe the anchor set rather than the literature.

Assignment simplification. Multi-method and multi-layer studies were assigned a single primary layer and method family.

Evidence base. About two-thirds of the core studies are preprints. Several anchor findings come from a small number of research groups and from developer-produced reports whose methods are not fully disclosed.

Recency. The field is moving quickly; the search date was 9 October 2026, and later studies are not included.

The most useful next steps are a human audit of a stratified sample of screening and charting decisions, full-text charting beyond the anchor set, and a registered living update of the review with the frozen query, extended to OpenReview and to hyphenated term variants.

8 Conclusion

Model-originated threats to evaluation validity have moved from theoretical concern to an active empirical field in less than three years. The evidence charted here shows that LLMs can recognise evaluations, that recognition can change behaviour, and that current detection methods are unreliable against adversarially constructed underperformance. It also shows that “evaluation awareness” is not one quantity, that its relationship to behaviour is model-specific, and that most validity evidence rests on behaviour that was deliberately induced. For practitioners, the immediate implications are to report awareness measures separately and with their biases, to prefer elicitation-first protocols for capability claims, and to treat propensity-based safety results as potentially biased in either direction when awareness is plausible, reporting cue-ablated or replay variants alongside them (Table 8, item 8). For researchers, the priority is to establish whether methods validated on induced behaviour transfer to behaviour that arises from training incentives.

Declarations

Funding The authors did not receive support from any organisation for the submitted work.

Competing interests The authors have no relevant financial or non-financial interests to disclose.

Ethics approval and consent to participate Not applicable. This study is a review of published and publicly available literature and involved no human participants, human data or animals.

Consent for publication Not applicable.

Data availability The protocol and its amendments, all search strategies and raw API responses, record logs, screening decisions, tier and charting data, full-text checks, reporting-audit codes and count reconciliations are provided as Supplementary Files S0–S5 and in the repository https://github.com/Bhargavteja-9779/AI_review_PNBT.git, which will be archived with a DOI on publication.

Code availability The scripts used for the database searches (`search/run_searches.py`), deduplication (`search/dedup.py`), screening and reconciliation (`screening/screen_db.py`, `screening/merge.py`, `screening/reliability.py`, `evidence/screen.py`), charting (`evidence/chart.py`) and figure generation (`evidence/figures.py`) are provided in the accompanying repository (https://github.com/Bhargavteja-9779/AI_review_PNBT.git), together with the raw API responses.

Author contributions P N Bhargav Teja: conceptualisation, methodology, investigation, data curation, formal analysis, visualisation, writing – original draft. Divya Meena S: conceptualisation, supervision, validation, writing – review and editing. Both authors read and approved the final manuscript.

Use of generative AI Generative AI tools, including Claude, were used as research assistance during the preparation of this review manuscript to support literature organization, thematic synthesis, academic language refinement, and manuscript structuring. All references, factual claims, and scientific interpretations were subject to human verification against the original sources. The authors take full responsibility for the accuracy, originality, integrity, and final content of the manuscript.

References

- Abdelnabi S, Salem A (2025) The Hawthorne effect in reasoning models: evaluating and steering test awareness. In: Advances in Neural Information Processing Systems (NeurIPS 2025). arXiv preprint arXiv:2505.14617. <https://arxiv.org/abs/2505.14617>
- Ahmed N, Zaman MI, Saleem G, Hassan A (2025) Do LLMs know they are being tested? Evaluation awareness and incentive-sensitive failures in GPT-OSS-20B. arXiv preprint arXiv:2510.08624. <https://arxiv.org/abs/2510.08624>
- Andriushchenko M, Souly A, Dziemian M et al. (2025) AgentHarm: a benchmark for measuring harmfulness of LLM agents. In: International Conference on Learning Representations (ICLR 2025). arXiv preprint arXiv:2410.09024. <https://arxiv.org/abs/2410.09024>
- Anthropic (2025a) System card: Claude Sonnet 4.5. <https://assets.anthropic.com/m/12f214efcc2f457a/original/Claude-Sonnet-4-5-System-Card.pdf>. Accessed 9 October 2026
- Anthropic (2025b) Responsible Scaling Policy, version 2.2. <https://www-cdn.anthropic.com/872c653b2d0501d6ab44cf87f43e1dc4853e4d37.pdf>. Accessed 9 October 2026
- Anthropic (2026) System card: Claude Opus 4.6. <https://anthropic.com/claude-opus-4-6-system-card>. Accessed 9 October 2026
- Anwar U, Abdelnabi S, Krueger D (2026) Training LLMs to verbalize evaluation awareness. arXiv preprint arXiv:2609.36316. <https://arxiv.org/abs/2609.36316>
- Arcuschin I, Janiak J, Krzyzanowski R, Rajamanoharan S, Nanda N, Conmy A (2026) Chain-of-thought reasoning in the wild is not always faithful. In: Proceedings of the 43rd International Conference on Machine Learning, PMLR 306. arXiv preprint arXiv:2503.08679. <https://arxiv.org/abs/2503.08679>
- Arditi A, Obeso O, Syed A et al. (2024) Refusal in language models is mediated by a single direction. In: Advances in Neural Information Processing Systems 37 (NeurIPS 2024). arXiv preprint arXiv:2406.11717. <https://arxiv.org/abs/2406.11717>
- Arksey H, O'Malley L (2005) Scoping studies: towards a methodological framework. Int J Soc Res Methodol 8(1):19–32. <https://doi.org/10.1080/1364557032000119616>
- Arnav B, Bernabeu-Pérez P, Helm-Burger N, Kostolansky T, Whittingham H, Phuong M (2025) CoT Red-Handed: stress testing chain-of-thought monitoring. In: Advances in Neural Information Processing Systems (NeurIPS 2025). arXiv preprint arXiv:2505.23575. <https://arxiv.org/abs/2505.23575>

Baek DD, Li X, Gupta A et al. (2026) Sycophancy towards researchers drives performative misalignment. In: Proceedings of the 43rd International Conference on Machine Learning, PMLR 306. arXiv preprint arXiv:2606.08629. <https://arxiv.org/abs/2606.08629>

Bailey L, Serrano A, Sheshadri A et al. (2026) Obfuscated activations bypass LLM latent-space defenses. In: International Conference on Learning Representations (ICLR 2026). arXiv preprint arXiv:2412.09565. <https://arxiv.org/abs/2412.09565>

Baker B, Huizinga J, Gao L et al. (2025) Monitoring reasoning models for misbehavior and the risks of promoting obfuscation. arXiv preprint arXiv:2503.11926. <https://arxiv.org/abs/2503.11926>

Balesni M, Hobbhahn M, Lindner D et al. (2024) Towards evaluations-based safety cases for AI scheming. arXiv preprint arXiv:2411.03336. <https://arxiv.org/abs/2411.03336>

Bean AM, Kearns RO, Romanou A et al. (2025) Measuring what matters: construct validity in large language model benchmarks. In: Advances in Neural Information Processing Systems (NeurIPS 2025). arXiv preprint arXiv:2511.04703. <https://arxiv.org/abs/2511.04703>

Benton J, Wagner M, Christiansen E et al. (2024) Sabotage evaluations for frontier models. arXiv preprint arXiv:2410.21514. <https://arxiv.org/abs/2410.21514>

Bercovich I, Segal I, Zhang K, Saxena S, Raghunathan A, Zhong Z (2026) Terminal Wrench: a dataset of 331 reward-hackable environments and 3,632 exploit trajectories. arXiv preprint arXiv:2604.17596. <https://arxiv.org/abs/2604.17596>

Berglund L, Stickland AC, Balesni M et al. (2023) Taken out of context: on measuring situational awareness in LLMs. arXiv preprint arXiv:2309.00667. <https://arxiv.org/abs/2309.00667>

Betley J, Tan D, Warncke N et al. (2025a) Emergent misalignment: narrow finetuning can produce broadly misaligned LLMs. In: Proceedings of the 42nd International Conference on Machine Learning, PMLR 267. arXiv preprint arXiv:2502.17424. <https://arxiv.org/abs/2502.17424>

Betley J, Bao X, Soto M, Szyber-Betley A, Chua J, Evans O (2025b) Tell me about yourself: LLMs are aware of their learned behaviors. In: International Conference on Learning Representations (ICLR 2025). arXiv preprint arXiv:2501.11120. <https://arxiv.org/abs/2501.11120>

Binder FJ, Chua J, Korbak T et al. (2025) Looking inward: language models can learn about themselves by introspection. In: International Conference on Learning Representations (ICLR 2025). arXiv preprint arXiv:2410.13787. <https://arxiv.org/abs/2410.13787>

Bondarenko A, Volk D, Volkov D, Ladish J (2026) Demonstrating specification gaming in reasoning models. In: TrustAgent Workshop at AAAI 2026. arXiv preprint arXiv:2502.13295. <https://arxiv.org/abs/2502.13295>

Bullwinkel B, Severi G, Hines K, Minnich A, Kumar RSS, Zunger Y (2026) The trigger in the haystack: extracting and reconstructing LLM backdoor triggers. arXiv preprint arXiv:2602.03085. <https://arxiv.org/abs/2602.03085>

Burnell R, Schellaert W, Burden J et al. (2023) Rethink reporting of evaluation results in AI. Science 380(6641):136–138. <https://doi.org/10.1126/science.adf6369>

Burns C, Ye H, Klein D, Steinhardt J (2023) Discovering latent knowledge in language models without supervision. In: International Conference on Learning Representations (ICLR 2023). arXiv preprint arXiv:2212.03827. <https://arxiv.org/abs/2212.03827>

Carlsmith J (2023) Scheming AIs: will AIs fake alignment during training in order to get power? arXiv preprint arXiv:2311.08379. <https://arxiv.org/abs/2311.08379>

Casper S, Ezell C, Siegmann C et al. (2024) Black-box access is insufficient for rigorous AI audits. In: Proceedings of the ACM Conference on Fairness, Accountability, and Transparency (FAccT 2024). arXiv preprint arXiv:2401.14446. <https://arxiv.org/abs/2401.14446>

Chaudhary M, Su I, Hooda N et al. (2025) Evaluation awareness scales predictably in open-weights large language models. In: ResponsibleFM Workshop at NeurIPS 2025. arXiv preprint arXiv:2509.13333. <https://arxiv.org/abs/2509.13333>

Chaudhary M (2026) In-context environments induce evaluation-awareness in language models. arXiv preprint arXiv:2603.03824. <https://arxiv.org/abs/2603.03824>

Che Z, Casper S, Kirk R et al. (2025) Model tampering attacks enable more rigorous evaluations of LLM capabilities. Transactions on Machine Learning Research. arXiv preprint arXiv:2502.05209. <https://arxiv.org/abs/2502.05209>

Chen B, Fang S, Ji J et al. (2025) AI deception: risks, dynamics, and controls. arXiv preprint arXiv:2511.22619. <https://arxiv.org/abs/2511.22619>

Chen S, Chen Y, Li Z et al. (2025) Benchmarking large language models under data contamination: a survey from static to dynamic evaluation. In: Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP 2025), pp 10091–10109. <https://doi.org/10.18653/v1/2025.emnlp-main.511>

Chen Y, Benton J, Radhakrishnan A et al. (2025) Reasoning models don't always say what they think. arXiv preprint arXiv:2505.05410. <https://arxiv.org/abs/2505.05410>

Chua J, Betley J, Taylor M, Evans O (2025) Thought crime: backdoors and emergent misalignment in reasoning models. arXiv preprint arXiv:2506.13206. <https://arxiv.org/abs/2506.13206>

Chua J, Evans O (2025) Are DeepSeek R1 and other reasoning models more faithful? In: ICLR 2025 Workshop on Foundation Models in the Wild. arXiv preprint arXiv:2501.08156. <https://arxiv.org/abs/2501.08156>

Clymer J, Juang C, Field S (2024) Poser: unmasking alignment faking LLMs by manipulating their internals. arXiv preprint arXiv:2405.05466. <https://arxiv.org/abs/2405.05466>

Cohen J (1960) A coefficient of agreement for nominal scales. *Educ Psychol Meas* 20(1):37–46. <https://doi.org/10.1177/001316446002000104>

Das A, Ponde J, More M, Gupta D (2026) You can't escape your own activations: evaluation awareness and multi-agent monitoring. arXiv preprint arXiv:2609.03035. <https://arxiv.org/abs/2609.03035>

Davidson T, Denain J-S, Villalobos P, Bas G (2023) AI capabilities can be significantly improved without expensive retraining. arXiv preprint arXiv:2312.07413. <https://arxiv.org/abs/2312.07413>

Denison C, MacDiarmid M, Barez F et al. (2024) Sycophancy to subterfuge: investigating reward-tampering in large language models. arXiv preprint arXiv:2406.10162. <https://arxiv.org/abs/2406.10162>

Devbunova V (2026) Is evaluation awareness just format sensitivity? limitations of probe-based evidence under controlled prompt structure. In: ICLR 2026 Workshop: I Can't Believe It's Not Better. arXiv preprint arXiv:2603.19426. <https://arxiv.org/abs/2603.19426>

- Dingeto H (2026) A lie detector test for language models: reading knowledge a model won't reveal. arXiv preprint arXiv:2609.21996. <https://arxiv.org/abs/2609.21996>
- Dubois Y, Galambosi B, Liang P, Hashimoto TB (2024) Length-controlled AlpacaEval: a simple way to debias automatic evaluators. In: Conference on Language Modeling (COLM 2024). arXiv preprint arXiv:2404.04475. <https://arxiv.org/abs/2404.04475>
- Emmons S, Jenner E, Elson DK et al. (2025) When chain of thought is necessary, language models struggle to evade monitors. arXiv preprint arXiv:2507.05246. <https://arxiv.org/abs/2507.05246>
- Fan Y, Zhang W, Pan X, Yang M (2025) Evaluation faking: unveiling observer effects in safety evaluation of frontier AI systems. arXiv preprint arXiv:2505.17815. <https://arxiv.org/abs/2505.17815>
- Ferrara E (2026) Defeat devices in AI systems. *Future Internet* 18(7):339. <https://doi.org/10.3390/fi18070339>
- Fonseca Rivera J, Africa DD (2026) Steering awareness: detecting activation steering from within. In: Conference on Language Modeling (COLM 2026). arXiv preprint arXiv:2511.21399. <https://arxiv.org/abs/2511.21399>
- Frontier Model Forum (2025) Frontier capability assessments. Technical report. <https://www.frontiermodelforum.org/technical-reports/frontier-capability-assessments/>. Accessed 9 October 2026
- Gao L, Schulman J, Hilton J (2023) Scaling laws for reward model overoptimization. In: Proceedings of the 40th International Conference on Machine Learning, PMLR 202:10835–10866. arXiv preprint arXiv:2210.10760. <https://arxiv.org/abs/2210.10760>
- Gardner-Challis N, Bostock J, Kozhevnikov G et al. (2026) When can we trust untrusted monitoring? a safety case sketch across collusion strategies. arXiv preprint arXiv:2602.20628. <https://arxiv.org/abs/2602.20628>
- Goldowsky-Dill N, Chughtai B, Heimersheim S, Hobbhahn M (2025) Detecting strategic deception using linear probes. In: Proceedings of the 42nd International Conference on Machine Learning (ICML 2025). arXiv preprint arXiv:2502.03407. <https://arxiv.org/abs/2502.03407>
- Greenblatt R, Denison C, Wright B et al. (2024a) Alignment faking in large language models. arXiv preprint arXiv:2412.14093. <https://arxiv.org/abs/2412.14093>
- Greenblatt R, Roger F, Krashennnikov D, Krueger D (2024b) Stress-testing capability elicitation with password-locked models. In: Advances in Neural Information Processing Systems 37 (NeurIPS 2024). arXiv preprint arXiv:2405.19550. <https://arxiv.org/abs/2405.19550>
- Greenblatt R, Shlegeris B, Sachan K, Roger F (2024c) AI control: improving safety despite intentional subversion. In: Proceedings of the 41st International Conference on Machine Learning, PMLR 235:16295–16336. arXiv preprint arXiv:2312.06942. <https://arxiv.org/abs/2312.06942>
- Grey M, Segerie C-R (2025) Safety by measurement: a systematic literature review of AI safety evaluation methods. arXiv preprint arXiv:2505.05541. <https://arxiv.org/abs/2505.05541>
- Gupta M, Nair I, Wang L, Kumar D (2026) Context over content: exposing evaluation faking in automated judges. arXiv preprint arXiv:2604.15224. <https://arxiv.org/abs/2604.15224>
- Hagendorff T (2024) Deception abilities emerged in large language models. *Proc Natl Acad Sci USA* 121(24):e2317967121. <https://doi.org/10.1073/pnas.2317967121>

- Hardy A, Reuel A, Meimandi KJ et al. (2025) More than marketing? on the information value of AI benchmarks for practitioners. In: Proceedings of the ACM Conference on Intelligent User Interfaces (IUI 2025). arXiv preprint arXiv:2412.05520. <https://arxiv.org/abs/2412.05520>
- Heidari F, Memarian A, Rabusseau G (2026) Evaluation awareness in language models: representation, verbalization, and control. arXiv preprint arXiv:2608.21766. <https://arxiv.org/abs/2608.21766>
- Hofstätter F, van der Weij T, Teoh J, Djoneva R, Bartsch H, Ward FR (2025) The elicitation game: evaluating capability elicitation techniques. In: Proceedings of the 42nd International Conference on Machine Learning, PMLR 267:23330–23356. arXiv preprint arXiv:2502.02180. <https://arxiv.org/abs/2502.02180>
- Hua TT, Qin A, Marks S, Nanda N (2026) Steering evaluation-aware language models to act like they are deployed. In: International Conference on Learning Representations (ICLR 2026). arXiv preprint arXiv:2510.20487. <https://arxiv.org/abs/2510.20487>
- Huang Y, Sun Y, Zhang Y, Zhang R, Dong Y, Wei X (2025) DeceptionBench: a comprehensive benchmark for AI deception behaviors in real-world scenarios. In: Advances in Neural Information Processing Systems (NeurIPS 2025), Datasets and Benchmarks Track. arXiv preprint arXiv:2510.15501. <https://arxiv.org/abs/2510.15501>
- Hubinger E, van Merwijk C, Mikulik V, Skalse J, Garrabrant S (2019) Risks from learned optimization in advanced machine learning systems. arXiv preprint arXiv:1906.01820. <https://arxiv.org/abs/1906.01820>
- Hubinger E, Denison C, Mu J et al. (2024) Sleeper agents: training deceptive LLMs that persist through safety training. arXiv preprint arXiv:2401.05566. <https://arxiv.org/abs/2401.05566>
- Ivanov I, Africa DD (2026) LURE: live-usage replay evaluations for reducing evaluation awareness. arXiv preprint arXiv:2605.26438. <https://arxiv.org/abs/2605.26438>
- Jacobs AZ, Wallach H (2021) Measurement and fairness. In: Proceedings of the ACM Conference on Fairness, Accountability, and Transparency (FAccT 2021). arXiv preprint arXiv:1912.05511. <https://arxiv.org/abs/1912.05511>
- Jang E, Falck D, Braun J et al. (2026) Exploration hacking: can LLMs learn to resist RL training? In: Proceedings of the 43rd International Conference on Machine Learning (ICML 2026). arXiv preprint arXiv:2604.28182. <https://arxiv.org/abs/2604.28182>
- Karvonen A, Chua J, Dumas C et al. (2026) Activation oracles: training and evaluating LLMs as general-purpose activation explainers. In: Proceedings of the 43rd International Conference on Machine Learning (ICML 2026). arXiv preprint arXiv:2512.15674. <https://arxiv.org/abs/2512.15674>
- Kaufman A, Lucassen J, Tracy T, Rushing C, Bhatt A (2025) BashArena: a control setting for highly privileged AI agents. arXiv preprint arXiv:2512.15688. <https://arxiv.org/abs/2512.15688>
- Kirk R, Souly A, Fronsdal K, D’Cruz A, Davies X (2026) Evaluating whether AI models would sabotage AI safety research. arXiv preprint arXiv:2604.24618. <https://arxiv.org/abs/2604.24618>
- Knecht A, Florin L, Hagendorff T (2026) Evaluation awareness in language models has limited effect on behaviour. arXiv preprint arXiv:2605.05835. <https://arxiv.org/abs/2605.05835>
- Korbak T, Balesni M, Barnes E et al. (2025) Chain of thought monitorability: a new and fragile opportunity for AI safety. arXiv preprint arXiv:2507.11473. <https://arxiv.org/abs/2507.11473>

- Kowalski MM, Fonseca Rivera J, Macar U, Africa DD (2026) Measuring activation control in large language models. arXiv preprint arXiv:2608.21664. <https://arxiv.org/abs/2608.21664>
- Krakovna V, Lindner D, Ho L, Farquhar S, Shah R (2026) Realistic honeypot evaluations for scheming propensity. arXiv preprint arXiv:2605.29729. <https://arxiv.org/abs/2605.29729>
- Kretschmar K, Laurito W, Maiya S, Marks S (2025) Liars’ Bench: evaluating lie detectors for language models. arXiv preprint arXiv:2511.16035. <https://arxiv.org/abs/2511.16035>
- Kumar S (2026) Pressure-testing deception probes in LLMs: scaling, robustness, and the geometry of deceptive representations. In: GEM Workshop at ACL 2026. arXiv preprint arXiv:2605.27958. <https://arxiv.org/abs/2605.27958>
- Kuntz T, Duzan A, Zhao H et al. (2025) OS-harm: a benchmark for measuring safety of computer use agents. In: Advances in Neural Information Processing Systems (NeurIPS 2025), Datasets and Benchmarks Track. arXiv preprint arXiv:2506.14866. <https://arxiv.org/abs/2506.14866>
- Kutasov J, Sun Y, Colognese P et al. (2025) SHADE-Arena: evaluating sabotage and monitoring in LLM agents. arXiv preprint arXiv:2506.15740. <https://arxiv.org/abs/2506.15740>
- Kwa T, West B, Becker J et al. (2025) Measuring AI ability to complete long software tasks. In: Advances in Neural Information Processing Systems (NeurIPS 2025). arXiv preprint arXiv:2503.14499. <https://arxiv.org/abs/2503.14499>
- Laine R, Chughtai B, Betley J et al. (2024) Me, myself, and AI: the situational awareness dataset (SAD) for LLMs. In: Advances in Neural Information Processing Systems 37 (NeurIPS 2024), Datasets and Benchmarks Track. arXiv preprint arXiv:2407.04694. <https://arxiv.org/abs/2407.04694>
- Lanham T, Chen A, Radhakrishnan A et al. (2023) Measuring faithfulness in chain-of-thought reasoning. arXiv preprint arXiv:2307.13702. <https://arxiv.org/abs/2307.13702>
- Le L, Tan HK, Williams-King D (2026) Reference-grafting matches fine-tuning at eliciting sandbagged capabilities. arXiv preprint arXiv:2608.29458. <https://arxiv.org/abs/2608.29458>
- Levac D, Colquhoun H, O’Brien KK (2010) Scoping studies: advancing the methodology. *Implement Sci* 5:69. <https://doi.org/10.1186/1748-5908-5-69>
- Li C, Phuong M, Siegel NY (2025) LLMs can covertly sandbag on capability evaluations against chain-of-thought monitoring. In: Proceedings of the 14th International Joint Conference on Natural Language Processing and 4th Conference of the Asia-Pacific Chapter of the ACL (IJCNLP-AAACL 2025). arXiv preprint arXiv:2508.00943. <https://arxiv.org/abs/2508.00943>
- Li C, Zhang TJ, Zhang J et al. (2026) Decomposing and measuring evaluation awareness. arXiv preprint arXiv:2605.23055. <https://arxiv.org/abs/2605.23055>
- Li JX, Turner A (2026) Eval cooperativeness mitigates evaluation gaming in LLMs. In: ICML 2026 Workshop on Trustworthy AI for Good. <https://icml.cc/virtual/2026/76891>
- Li N, Pan A, Gopal A et al. (2024) The WMDP benchmark: measuring and reducing malicious use with unlearning. In: Proceedings of the 41st International Conference on Machine Learning, PMLR 235:28525–28550. arXiv preprint arXiv:2403.03218. <https://arxiv.org/abs/2403.03218>
- Li X, Ochwang’i K, Bharadwaj AR, Souly A, Kirk R (2026) EvalDetectBench: a benchmark for measuring evaluation awareness in frontier language models. arXiv preprint arXiv:2609.01611.

<https://arxiv.org/abs/2609.01611>

Liang P, Bommasani R, Lee T et al. (2023) Holistic evaluation of language models. Transactions on Machine Learning Research. arXiv preprint arXiv:2211.09110. <https://arxiv.org/abs/2211.09110>

Lin Y, Bernabeu-Perez P, Arnav B, Wells L, Phuong M (2026a) Same question, different lies: cross-context consistency (C³) for black-box sandbagging detection. In: Proceedings of the 43rd International Conference on Machine Learning (ICML 2026); earlier version at the ICLR 2026 Workshop on Principled Design for Trustworthy AI. <https://openreview.net/forum?id=ldgh8zhujD>

Lin Y, Bernabeu-Perez P, Arnav B, Wells L, Phuong M (2026b) When lie detectors learn model identity: confounds in black-box sandbagging detection. In: ICLR 2026 Workshop: I Can't Believe It's Not Better. <https://openreview.net/forum?id=rb7BWE6qFE>

Łucki J, Wei B, Huang Y, Henderson P, Tramèr F, Rando J (2025) An adversarial perspective on machine unlearning for AI safety. Transactions on Machine Learning Research. arXiv preprint arXiv:2409.18025. <https://arxiv.org/abs/2409.18025>

Lynch A, Wright B, Larson C et al. (2025) Agentic misalignment: how LLMs could be insider threats. arXiv preprint arXiv:2510.05179. <https://arxiv.org/abs/2510.05179>

Mallen A, Brumley M, Kharchenko J, Belrose N (2024) Eliciting latent knowledge from quirky language models. In: Conference on Language Modeling (COLM 2024). arXiv preprint arXiv:2312.01037. <https://arxiv.org/abs/2312.01037>

Manek A (2026) Representational depth of evaluation awareness shifts with scale in open-weight language models. In: ICML 2026 Workshop on Mechanistic Interpretability. arXiv preprint arXiv:2606.29196. <https://arxiv.org/abs/2606.29196>

Marks S, Tegmark M (2024) The geometry of truth: emergent linear structure in large language model representations of true/false datasets. In: Conference on Language Modeling (COLM 2024). arXiv preprint arXiv:2310.06824. <https://arxiv.org/abs/2310.06824>

Marks S, Treutlein J, Bricken T et al. (2025) Auditing language models for hidden objectives. arXiv preprint arXiv:2503.10965. <https://arxiv.org/abs/2503.10965>

Mathew Y, Matthews O, McCarthy R et al. (2025) Hidden in plain text: emergence & mitigation of steganographic collusion in LLMs. In: Proceedings of the 14th International Joint Conference on Natural Language Processing and 4th Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics (IJCNLP-AACL 2025). arXiv preprint arXiv:2410.03768. <https://arxiv.org/abs/2410.03768>

Maurya J, Chaurasiya S, Siddiqui F, Perwej Y (2026) A systematic review of evaluation of how AI systems behaves when unmonitored. Int J Sci Res Comput Sci Eng Inf Technol 12(2):642–667. <https://doi.org/10.32628/CSEIT26121384>

Mazeika M, Phan L, Yin X et al. (2024) HarmBench: a standardized evaluation framework for automated red teaming and robust refusal. In: Proceedings of the 41st International Conference on Machine Learning, PMLR 235. arXiv preprint arXiv:2402.04249. <https://arxiv.org/abs/2402.04249>

McCambridge J, Witton J, Elbourne DR (2014) Systematic review of the Hawthorne effect: new concepts are needed to study research participation effects. J Clin Epidemiol 67(3):267–277. <https://doi.org/10.1016/j.jclinepi.2013.08.015>

- Mehta A (2026) The refusal residue: when probes catch alignment faking and when they don't. In: ICML 2026 Workshop on Mechanistic Interpretability. arXiv preprint arXiv:2607.13346. <https://arxiv.org/abs/2607.13346>
- Meinke A, Schoen B, Scheurer J, Balesni M, Shah R, Hobbhahn M (2024) Frontier models are capable of in-context scheming. arXiv preprint arXiv:2412.04984. <https://arxiv.org/abs/2412.04984>
- Messick S (1995) Validity of psychological assessment: validation of inferences from persons' responses and performances as scientific inquiry into score meaning. *Am Psychol* 50(9):741–749. <https://doi.org/10.1037/0003-066X.50.9.741>
- METR (2025a) Details about METR's evaluation of OpenAI GPT-5. <https://metr.org/evaluations/gpt-5-report/>. Accessed 9 October 2026
- METR (2025b) Details about METR's preliminary evaluation of OpenAI's o3 and o4-mini. <https://metr.org/evaluations/openai-o3-report/>. Accessed 9 October 2026
- Mizrahi M, Kaplan G, Malkin D, Dror R, Shahaf D, Stanovsky G (2024) State of what art? a call for multi-prompt LLM evaluation. *Transactions of the Association for Computational Linguistics* 12. arXiv preprint arXiv:2401.00595. <https://arxiv.org/abs/2401.00595>
- Mody D, Agarwal S, Mittal U, Mahato D (2026) Minimizing targeted activations: input-only suppression of evaluation-awareness latents in large language models. arXiv preprint arXiv:2607.25907. <https://arxiv.org/abs/2607.25907>
- Montoya LM, Haskins R, Adams B (2026) Towards measuring and detecting un verbalized evaluation awareness. In: ICML 2026 Workshop on Mechanistic Interpretability. <https://openreview.net/forum?id=zGmgeR0f4D>
- Morampudi A, Irrinki U, Grandhi R, Pagadala V, Maddula M (2026) A survey of reward hacking in agentic large language model systems. *Discov Artif Intell* 6:825. <https://doi.org/10.1007/s44163-026-01980-z>
- Motwani SR, Baranchuk M, Strohmeier M et al. (2024) Secret collusion among AI agents: multi-agent deception via steganography. In: *Advances in Neural Information Processing Systems* 37 (NeurIPS 2024). arXiv preprint arXiv:2402.07510. <https://arxiv.org/abs/2402.07510>
- Nayan N, Sampath Kumar A, Girmal R et al. (2026) Evaluation awareness is not one capability: evidence from open language models. arXiv preprint arXiv:2606.23583. <https://arxiv.org/abs/2606.23583>
- Needham J, Edkins G, Pimpale G, Bartsch H, Hobbhahn M (2025) Large language models often know when they are being evaluated. arXiv preprint arXiv:2505.23836. <https://arxiv.org/abs/2505.23836>
- Ngo R, Chan L, Mindermann S (2024) The alignment problem from a deep learning perspective. In: *International Conference on Learning Representations (ICLR 2024)*. arXiv preprint arXiv:2209.00626. <https://arxiv.org/abs/2209.00626>
- Nguyen J, Hoang K, Attubato CL, Hofstätter F (2025) Probing and steering evaluation awareness of language models. In: *ICML 2025 Workshop on Actionable Interpretability*. arXiv preprint arXiv:2507.01786. <https://arxiv.org/abs/2507.01786>
- Noël V (2026) A probe direction is a property of its prompt. arXiv preprint arXiv:2608.13329. <https://arxiv.org/abs/2608.13329>
- OpenAI (2025a) GPT-5 system card. arXiv preprint arXiv:2601.03267. <https://arxiv.org/abs/2601.03267>

OpenAI (2025b) Preparedness Framework, version 2. <https://openai.com/index/updating-our-preparedness-framework/>. Accessed 9 October 2026

Oren Y, Meister N, Chatterji N, Ladhak F, Hashimoto TB (2024) Proving test set contamination in black box language models. In: International Conference on Learning Representations (ICLR 2024). arXiv preprint arXiv:2310.17623. <https://arxiv.org/abs/2310.17623>

Pacchiardi L, Chan AJ, Mindermann S et al. (2024) How to catch an AI liar: lie detection in black-box LLMs by asking unrelated questions. In: International Conference on Learning Representations (ICLR 2024). arXiv preprint arXiv:2309.15840. <https://arxiv.org/abs/2309.15840>

Page MJ, McKenzie JE, Bossuyt PM et al. (2021) The PRISMA 2020 statement: an updated guideline for reporting systematic reviews. *BMJ* 372:n71. <https://doi.org/10.1136/bmj.n71>

Pan X, Dai J, Fan Y, Yang M (2024) Frontier AI systems have surpassed the self-replicating red line. arXiv preprint arXiv:2412.12140. <https://arxiv.org/abs/2412.12140>

Panfilov A, Kortukov E, Nikolić K et al. (2026) Strategic dishonesty can undermine AI safety evaluations of frontier LLMs. In: International Conference on Learning Representations (ICLR 2026). arXiv preprint arXiv:2509.18058. <https://arxiv.org/abs/2509.18058>

Panickssery A, Bowman SR, Feng S (2024) LLM evaluators recognize and favor their own generations. In: Advances in Neural Information Processing Systems 37 (NeurIPS 2024). arXiv preprint arXiv:2404.13076. <https://arxiv.org/abs/2404.13076>

Park PS, Goldstein S, O’Gara A, Chen M, Hendrycks D (2024) AI deception: a survey of examples, risks, and potential solutions. *Patterns* 5(5):100988. <https://doi.org/10.1016/j.patter.2024.100988>

Parrack A, Attubato CL, Heimersheim S (2025) Benchmarking deception probes via black-to-white performance boosts. arXiv preprint arXiv:2507.12691. <https://arxiv.org/abs/2507.12691>

Phuong M, Aitchison M, Catt E et al. (2024) Evaluating frontier models for dangerous capabilities. arXiv preprint arXiv:2403.13793. <https://arxiv.org/abs/2403.13793>

Phuong M, Zimmermann RS, Wang Z et al. (2025) Evaluating frontier models for stealth and situational awareness. arXiv preprint arXiv:2505.01420. <https://arxiv.org/abs/2505.01420>

Price S, Panickssery A, Bowman S, Stickland AC (2024) Future events as backdoor triggers: investigating temporal vulnerabilities in LLMs. arXiv preprint arXiv:2407.04108. <https://arxiv.org/abs/2407.04108>

Qin J, Schmotz D, Prinzhorn D, Beurer-Kellner L, Prabhu A, Andriushchenko M (2026) LLM agents can easily tamper with their own traces. arXiv preprint arXiv:2609.30266. <https://arxiv.org/abs/2609.30266>

Rajan S (2026) Auditing reward hackability in code RL training environments. arXiv preprint arXiv:2606.16062. <https://arxiv.org/abs/2606.16062>

Raji ID, Bender EM, Paullada A, Denton E, Hanna A (2021) AI and the everything in the whole wide world benchmark. In: NeurIPS 2021 Datasets and Benchmarks Track. arXiv preprint arXiv:2111.15366. <https://arxiv.org/abs/2111.15366>

Ren R, Agarwal A, Mazeika M et al. (2025) The MASK benchmark: disentangling honesty from accuracy in AI systems. arXiv preprint arXiv:2503.03750. <https://arxiv.org/abs/2503.03750>

- Reuel A, Hardy A, Smith C, Lamparth M, Hardy M, Kochenderfer MJ (2024) BetterBench: assessing AI benchmarks, uncovering issues, and establishing best practices. In: Advances in Neural Information Processing Systems 37 (NeurIPS 2024), Datasets and Benchmarks Track. arXiv preprint arXiv:2411.12990. <https://arxiv.org/abs/2411.12990>
- Rimsky N, Gabrieli N, Schulz J, Tong M, Hubinger E, Turner AM (2024) Steering Llama 2 via contrastive activation addition. In: Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), pp 15504–15522. arXiv preprint arXiv:2312.06681. <https://arxiv.org/abs/2312.06681>
- Roger F, Greenblatt R (2023) Preventing language models from hiding their reasoning. arXiv preprint arXiv:2310.18512. <https://arxiv.org/abs/2310.18512>
- Rosu P, Wang R (2026) Training alignment auditors via reinforcement learning. arXiv preprint arXiv:2608.25460. <https://arxiv.org/abs/2608.25460>
- Röttger P, Pernisi F, Vidgen B, Hovy D (2025) SafetyPrompts: a systematic review of open datasets for evaluating and improving large language model safety. In: Proceedings of the AAAI Conference on Artificial Intelligence (AAAI 2025). arXiv preprint arXiv:2404.05399. <https://arxiv.org/abs/2404.05399>
- Rozinek O (2026) How many queries to catch a sandbagger? A Chernoff–Stein sample-complexity theory of capability elicitation. In: Proceedings of the 23rd International Conference on Security and Cryptography. SCITEPRESS, pp 383–390. <https://doi.org/10.5220/0015190400004103>
- Ryd E, Bartsch H, Stastny J, Benton J, Hebbar V (2026) Removing sandbagging in LLMs by training with weak supervision. In: Proceedings of the 43rd International Conference on Machine Learning, PMLR 306. arXiv preprint arXiv:2604.22082. <https://arxiv.org/abs/2604.22082>
- Santos-Grueiro I (2026) Alignment verifiability in large language models: normative indistinguishability under behavioral evaluation. arXiv preprint arXiv:2602.05656. <https://arxiv.org/abs/2602.05656>
- Saraireh O (2026) Elicitation floors: what a capability evaluation cannot report when refusals are excluded. SSRN preprint 7402158. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=7402158. Accessed 9 October 2026
- Scheurer J, Balesni M, Hobbhahn M (2024) Large language models can strategically deceive their users when put under pressure. In: ICLR 2024 Workshop on Large Language Model (LLM) Agents. arXiv preprint arXiv:2311.07590. <https://arxiv.org/abs/2311.07590>
- Schlatter J, Weinstein-Raun B, Ladish J (2025) Incomplete tasks induce shutdown resistance in some frontier LLMs. Transactions on Machine Learning Research. arXiv preprint arXiv:2509.14260. <https://arxiv.org/abs/2509.14260>
- Schoen B, Nitishinskaya E, Balesni M et al. (2025) Stress testing deliberative alignment for anti-scheming training. arXiv preprint arXiv:2509.15541. <https://arxiv.org/abs/2509.15541>
- Sclar M, Choi Y, Tsvetkov Y, Suhr A (2024) Quantifying language models’ sensitivity to spurious features in prompt design or: how I learned to start worrying about prompt formatting. In: International Conference on Learning Representations (ICLR 2024). arXiv preprint arXiv:2310.11324. <https://arxiv.org/abs/2310.11324>
- Sheshadri A, Hughes J, Michael J et al. (2025) Why do some language models fake alignment while others don’t? In: Advances in Neural Information Processing Systems (NeurIPS 2025). arXiv preprint arXiv:2506.18032. <https://arxiv.org/abs/2506.18032>

- Shi J, Zhang TJ, Jin Z, Conitzer V (2026) From sycophancy to deception: a unified taxonomy for LLM spontaneous misalignment. In: ICLR 2026 Workshop: Agents in the Wild: Safety, Security, and Beyond. arXiv preprint arXiv:2604.04788. <https://arxiv.org/abs/2604.04788>
- Singh N, Chaudhary M (2026) Evaluation awareness shifts from format to context with model scale. arXiv preprint arXiv:2609.22119. <https://arxiv.org/abs/2609.22119>
- Singh S, Nan Y, Wang A et al. (2025) The leaderboard illusion. In: Advances in Neural Information Processing Systems (NeurIPS 2025). arXiv preprint arXiv:2504.20879. <https://arxiv.org/abs/2504.20879>
- Soligo A, Turner E, Rajamanoharan S, Nanda N (2025) Convergent linear representations of emergent misalignment. In: NeurIPS 2025 Workshop on Mechanistic Interpretability. arXiv preprint arXiv:2506.11618. <https://arxiv.org/abs/2506.11618>
- Souly A, Lu Q, Bowen D et al. (2024) A StrongREJECT for empty jailbreaks. In: Advances in Neural Information Processing Systems 37 (NeurIPS 2024), Datasets and Benchmarks Track. arXiv preprint arXiv:2402.10260. <https://arxiv.org/abs/2402.10260>
- Souly A, Kirk R, Merizian J, D’Cruz A, Davies X (2026) UK AISI alignment evaluation case-study. arXiv preprint arXiv:2604.00788. <https://arxiv.org/abs/2604.00788>
- Tamirisa R, Bharathi B, Phan L et al. (2025) Tamper-resistant safeguards for open-weight LLMs. In: International Conference on Learning Representations (ICLR 2025). arXiv preprint arXiv:2408.00761. <https://arxiv.org/abs/2408.00761>
- Tan D, Woodruff A, Warncke N et al. (2025) Inoculation prompting: eliciting traits from LLMs during training can suppress them at test-time. arXiv preprint arXiv:2510.04340. <https://arxiv.org/abs/2510.04340>
- Taylor J, Black S, Bowen D et al. (2025) Auditing games for sandbagging. arXiv preprint arXiv:2512.07810. <https://arxiv.org/abs/2512.07810>
- Taylor M, Chua J, Betley J, Treutlein J, Evans O (2025) School of reward hacks: hacking harmless tasks generalizes to misaligned behavior in LLMs. arXiv preprint arXiv:2508.17511. <https://arxiv.org/abs/2508.17511>
- Tice C, Kreer PA, Helm-Burger N et al. (2025) Noise injection reveals hidden capabilities of sandbagging language models. In: Advances in Neural Information Processing Systems (NeurIPS 2025). arXiv preprint arXiv:2412.01784. <https://arxiv.org/abs/2412.01784>
- Treutlein J, Choi D, Betley J et al. (2024) Connecting the dots: LLMs can infer and verbalize latent structure from disparate training data. In: Advances in Neural Information Processing Systems 37 (NeurIPS 2024). arXiv preprint arXiv:2406.14546. <https://arxiv.org/abs/2406.14546>
- Tricco AC, Lillie E, Zarin W et al. (2018) PRISMA extension for scoping reviews (PRISMA-ScR): checklist and explanation. *Ann Intern Med* 169(7):467–473. <https://doi.org/10.7326/M18-0850>
- Turner E, Soligo A, Taylor M, Rajamanoharan S, Nanda N (2025) Model organisms for emergent misalignment. arXiv preprint arXiv:2506.11613. <https://arxiv.org/abs/2506.11613>
- Turpin M, Michael J, Perez E, Bowman SR (2023) Language models don’t always say what they think: unfaithful explanations in chain-of-thought prompting. In: Advances in Neural Information Processing Systems 36 (NeurIPS 2023). arXiv preprint arXiv:2305.04388. <https://arxiv.org/abs/2305.04388>

- Uluirmak BA, Kurban R (2026) EvalSafetyGap: a hybrid survey and conceptual framework for LLM evaluation-safety failures. arXiv preprint arXiv:2606.30219. <https://arxiv.org/abs/2606.30219>
- Vu HM, Michala J, Saunders W (2026) Toward dealing with un verbalized eval awareness. In: ICML 2026 Workshop on Trustworthy AI for Good (AI4GOOD). <https://openreview.net/forum?id=VgHY8vLww3>
- Wang H, Li H, Mang Q, Cheung A, Sen K, Song D (2026) Do androids dream of breaking the game? Systematically auditing AI agent benchmarks with BenchJack. arXiv preprint arXiv:2605.12673. <https://arxiv.org/abs/2605.12673>
- Wang M, Dupré la Tour T, Watkins O et al. (2026) Persona features control emergent misalignment. In: International Conference on Learning Representations (ICLR 2026). arXiv preprint arXiv:2506.19823. <https://arxiv.org/abs/2506.19823>
- Wang P, Li L, Chen L et al. (2024) Large language models are not fair evaluators. In: Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), pp 9440–9450. arXiv preprint arXiv:2305.17926. <https://arxiv.org/abs/2305.17926>
- Wang X, Tian M, Zeng Y et al. (2026) Reward hacking in the era of large models: mechanisms, emergent misalignment, challenges. arXiv preprint arXiv:2604.13602. <https://arxiv.org/abs/2604.13602>
- Ward FR, van der Weij T, Gábor H et al. (2025) CTRL-ALT-DECEIT: sabotage evaluations for automated AI R&D. In: Advances in Neural Information Processing Systems (NeurIPS 2025). arXiv preprint arXiv:2511.09904. <https://arxiv.org/abs/2511.09904>
- Weidinger L, Raji ID, Wallach H et al. (2025) Toward an evaluation science for generative AI systems. arXiv preprint arXiv:2503.05336. <https://arxiv.org/abs/2503.05336>
- van der Weij T, Hofstätter F, Jaffe O, Brown SF, Ward FR (2025) AI sandbagging: language models can strategically underperform on evaluations. In: International Conference on Learning Representations (ICLR 2025). arXiv preprint arXiv:2406.07358. <https://arxiv.org/abs/2406.07358>
- Wichers N, Ebtekar A, Azarbal A et al. (2025) Inoculation prompting: instructing LLMs to misbehave at train-time improves test-time alignment. arXiv preprint arXiv:2510.05024. <https://arxiv.org/abs/2510.05024>
- Williams M, Sheahan H, Raymond C et al. (2026) Predicting LLM safety before release by simulating deployment. arXiv preprint arXiv:2607.07184. <https://arxiv.org/abs/2607.07184>
- Wu Y, Gao Q, Pan X, Hong G, Yang M (2026) OpenDeception: learning deception and trust in Human-AI interaction via multi-agent simulation. In: Proceedings of the 43rd International Conference on Machine Learning (ICML 2026). arXiv preprint arXiv:2504.13707. <https://arxiv.org/abs/2504.13707>
- Xiong L, Bhargava N, Hong J et al. (2025) Probe-Rewrite-Evaluate: a workflow for reliable benchmarks and quantifying evaluation awareness. arXiv preprint arXiv:2509.00591. <https://arxiv.org/abs/2509.00591>
- Xu X (2026) The oversight gap: what LLM safety monitors miss, and why it is not capability. arXiv preprint arXiv:2609.07162. <https://arxiv.org/abs/2609.07162>
- Zhong Z, Raghunathan A, Carlini N (2026a) ImpossibleBench: measuring LLMs’ propensity of exploiting test cases. In: International Conference on Learning Representations (ICLR 2026). arXiv preprint arXiv:2510.20270. <https://arxiv.org/abs/2510.20270>
- Zhong Z, Segal I, Bercovich I, Saxena S, Zhang K, Raghunathan A (2026b) Hardening agent benchmarks with adversarial hacker-fixer loops. arXiv preprint arXiv:2606.08960. <https://arxiv.org/abs/2606.08960>

Zhu Y, Jin T, Pruksachatkun Y et al. (2025) Establishing best practices for building rigorous agentic benchmarks. In: Advances in Neural Information Processing Systems (NeurIPS 2025), Datasets and Benchmarks Track. arXiv preprint arXiv:2507.02825. <https://arxiv.org/abs/2507.02825>

Zhuang A, Aranguri S (2026) Not all eval-awareness is equal: capabilities framing predicts compliance. In: ICML 2026 Workshop on Mechanistic Interpretability. arXiv preprint arXiv:2608.27340. <https://arxiv.org/abs/2608.27340>

Zou A, Phan L, Chen S et al. (2023) Representation engineering: a top-down approach to AI transparency. arXiv preprint arXiv:2310.01405. <https://arxiv.org/abs/2310.01405>